

# Preparation

Please go to <https://github.com/ray-project/tutorial>

And click on the “Binder” link and wait a couple minutes.

## Try Ray on Binder (Experimental)

---

Try the Ray tutorials online on [Binder](#).

Binder not working? Backup:

- run locally with "pip install ray[rllib] tensorflow jupyterlab"
- sign in to Google CoLab and import the tutorial notebooks



#ODSC

**BOSTON**  
APR 30 – MAY 3

# Scaling AI Applications with Ray

A general-purpose system for parallel and distributed Python

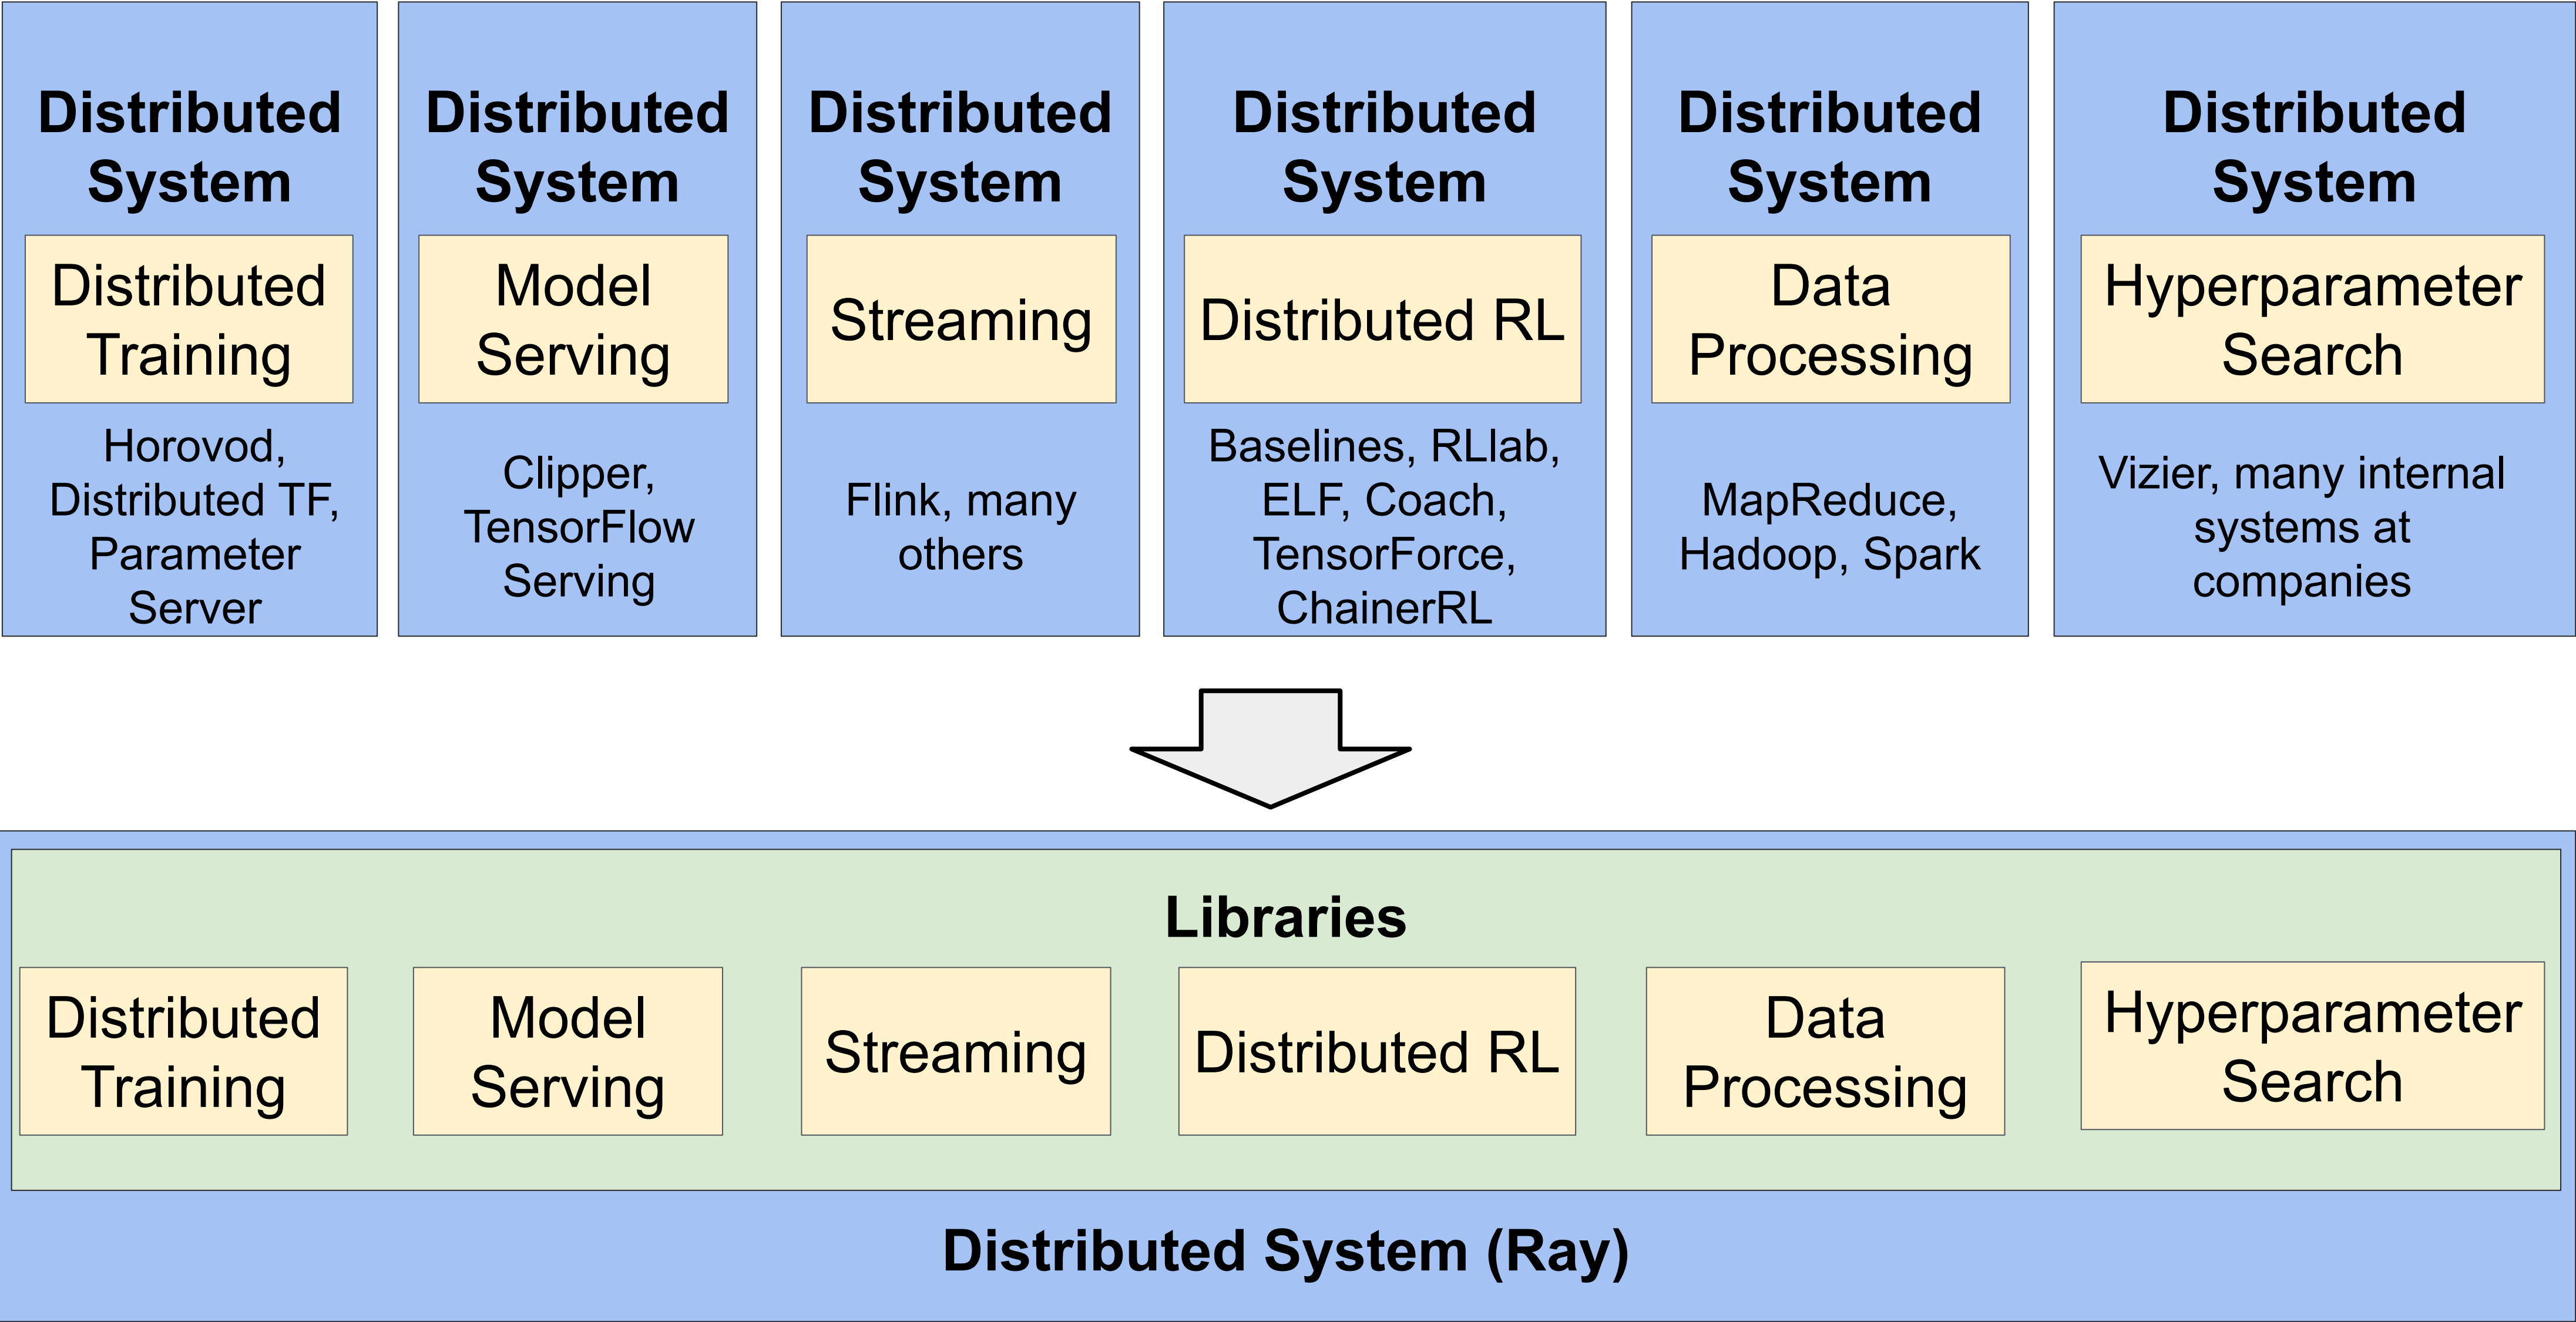
<https://github.com/ray-project/ray>

**Richard Liaw**

AI Researcher,  
RISELab at UC Berkeley




# The Big Picture





# A Growing Number of Use Cases



J.P.Morgan

Morgan Stanley

PRIMER



# Agenda

**Part 1 (2:00-2:30):** Ray as a unifying distributed system for machine learning

- Ray API, Architecture
- Exercises

**Part 2 (2:30-3:00):** Model Tuning/Training

- Overview of Tune
- Tune Exercise

**Part 3 (3:00-3:30):** Applied Reinforcement Learning

- Overview of RLlib
- RLlib Exercise



# Ray API

## Functions -> Tasks

```
def read_array(file):  
    # read array “a” from “file”  
    return a
```

```
def add(a, b):  
    return np.add(a, b)
```



# Ray API

## Functions -> Tasks

```
@ray.remote
```

```
def read_array(file):  
    # read array “a” from “file”  
    return a
```

```
@ray.remote
```

```
def add(a, b):  
    return np.add(a, b)
```




# Ray API

## Functions -> Tasks

```
@ray.remote
def read_array(file):
    # read array "a" from "file"
    return a
```

```
@ray.remote
def add(a, b):
    return np.add(a, b)
```

```
id1 = read_array.remote([5, 5])
id2 = read_array.remote([5, 5])
```





# Ray API

## Functions -> Tasks


```
@ray.remote
```

```
def read_array(file):  
    # read array "a" from "file"  
    return a
```

```
@ray.remote
```

```
def add(a, b):  
    return np.add(a, b)
```

```
id1 = read_array.remote([5, 5])  
id2 = read_array.remote([5, 5])  
id3 = add.remote(id1, id2)
```



# Ray API

## Functions -> Tasks


```
@ray.remote
```

```
def read_array(file):  
    # read array "a" from "file"  
    return a
```

```
@ray.remote
```

```
def add(a, b):  
    return np.add(a, b)
```

```
id1 = read_array.remote([5, 5])  
id2 = read_array.remote([5, 5])  
id3 = add.remote(id1, id2)  
ray.get(id3)
```



# Ray API

## Functions -> Tasks

```
@ray.remote
def read_array(file):
    # read array "a" from "file"
    return a
```

```
@ray.remote
def add(a, b):
    return np.add(a, b)
```

```
id1 = read_array.remote([5, 5])
id2 = read_array.remote([5, 5])
id3 = add.remote(id1, id2)
ray.get(id3)
```

## Classes -> Actors



# Ray API

## Functions -> Tasks

```
@ray.remote
def read_array(file):
    # read array "a" from "file"
    return a
```

```
@ray.remote
def add(a, b):
    return np.add(a, b)
```

```
id1 = read_array.remote([5, 5])
id2 = read_array.remote([5, 5])
id3 = add.remote(id1, id2)
ray.get(id3)
```

## Classes -> Actors

```
@ray.remote
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
    return self.value
```



# Ray API

## Functions -> Tasks

```
@ray.remote
def read_array(file):
    # read array "a" from "file"
    return a
```

```
@ray.remote
def add(a, b):
    return np.add(a, b)
```

```
id1 = read_array.remote([5, 5])
id2 = read_array.remote([5, 5])
id3 = add.remote(id1, id2)
ray.get(id3)
```

## Classes -> Actors

```
@ray.remote
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
    return self.value
```

```
c = Counter.remote()
id4 = c.inc.remote()
id5 = c.inc.remote()
ray.get([id4, id5])
```

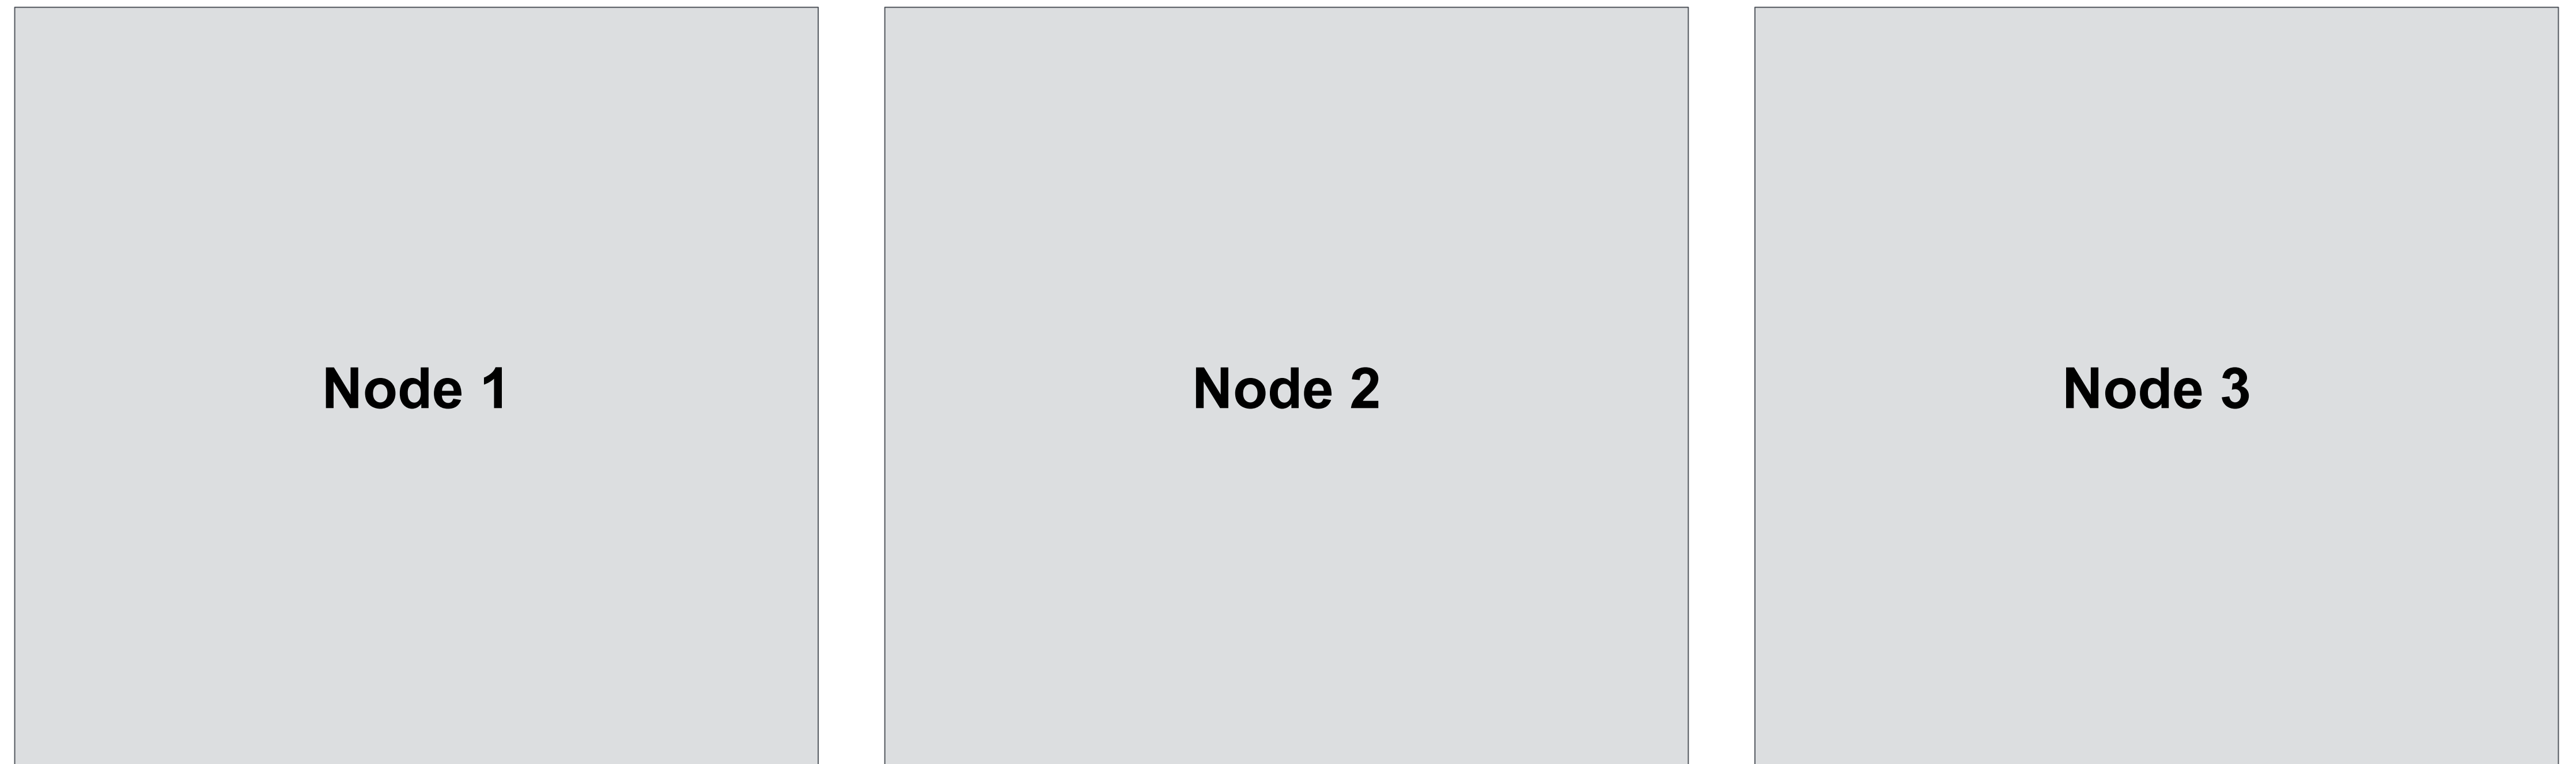




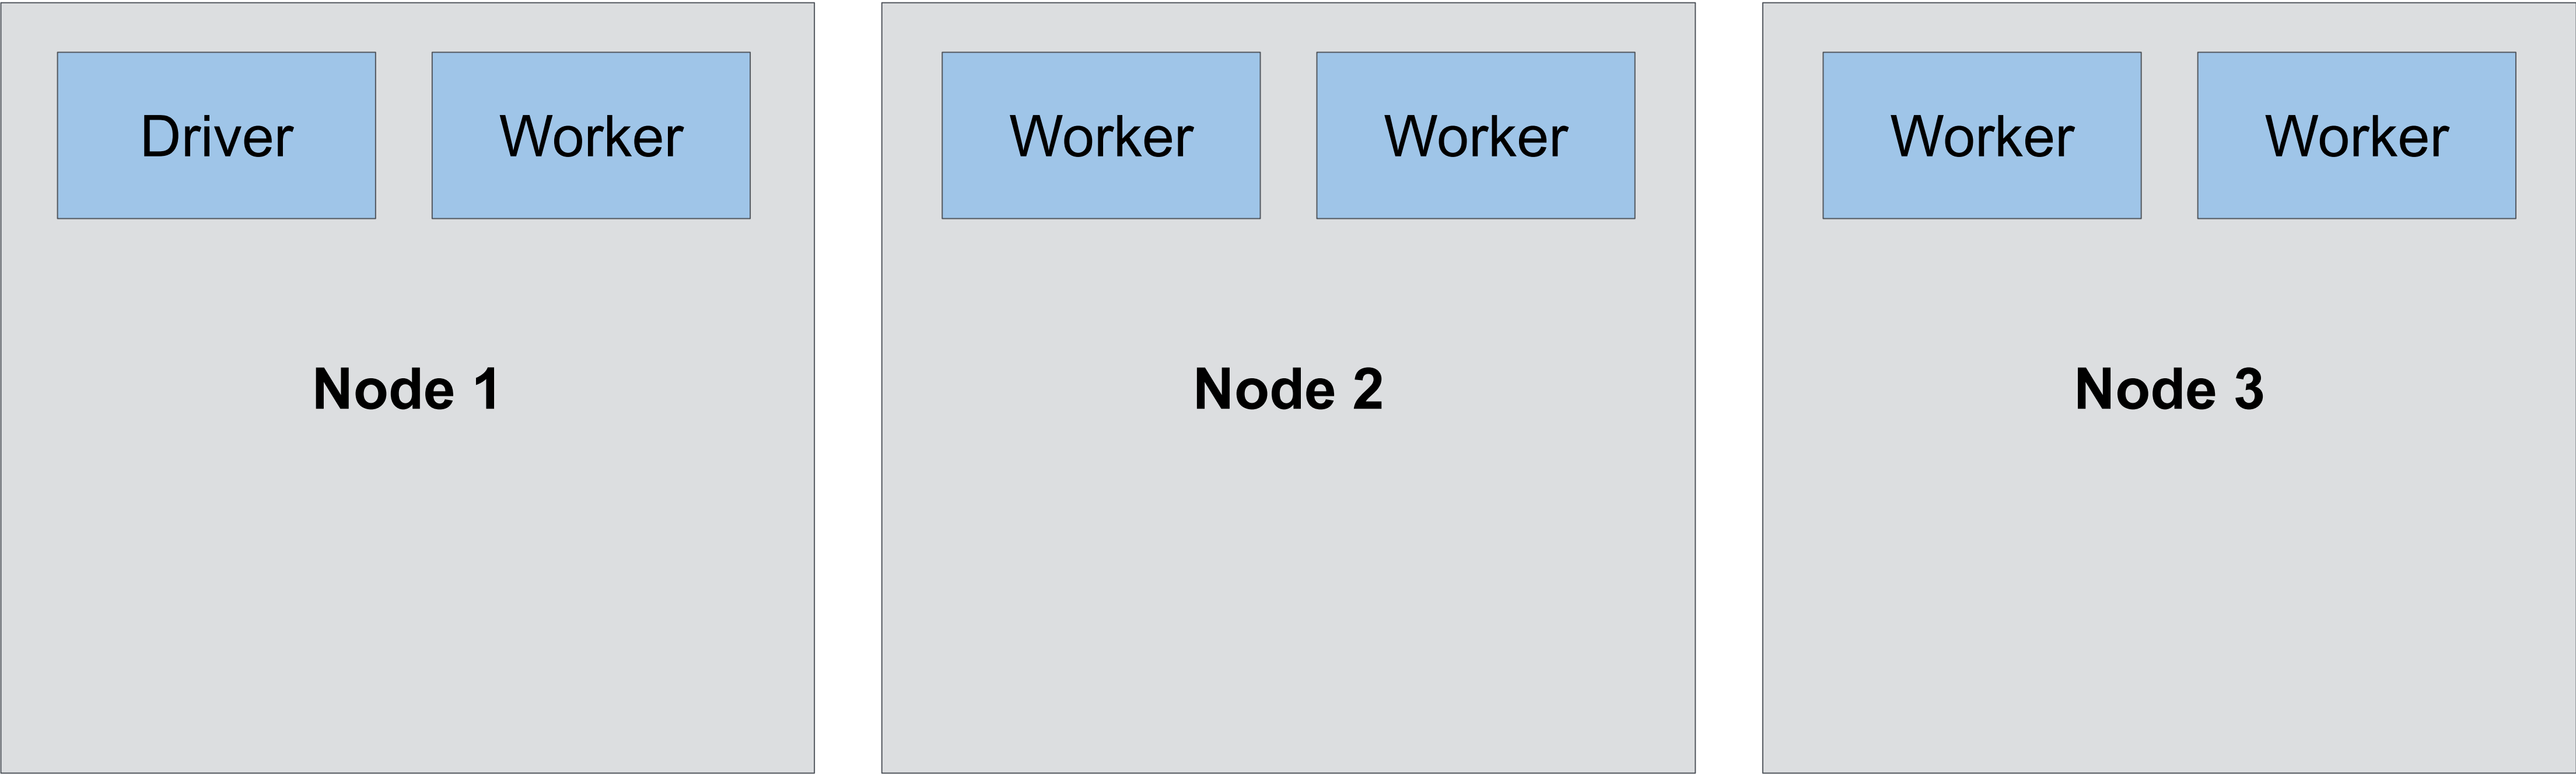


# Quick Overview of **Ray Architecture**

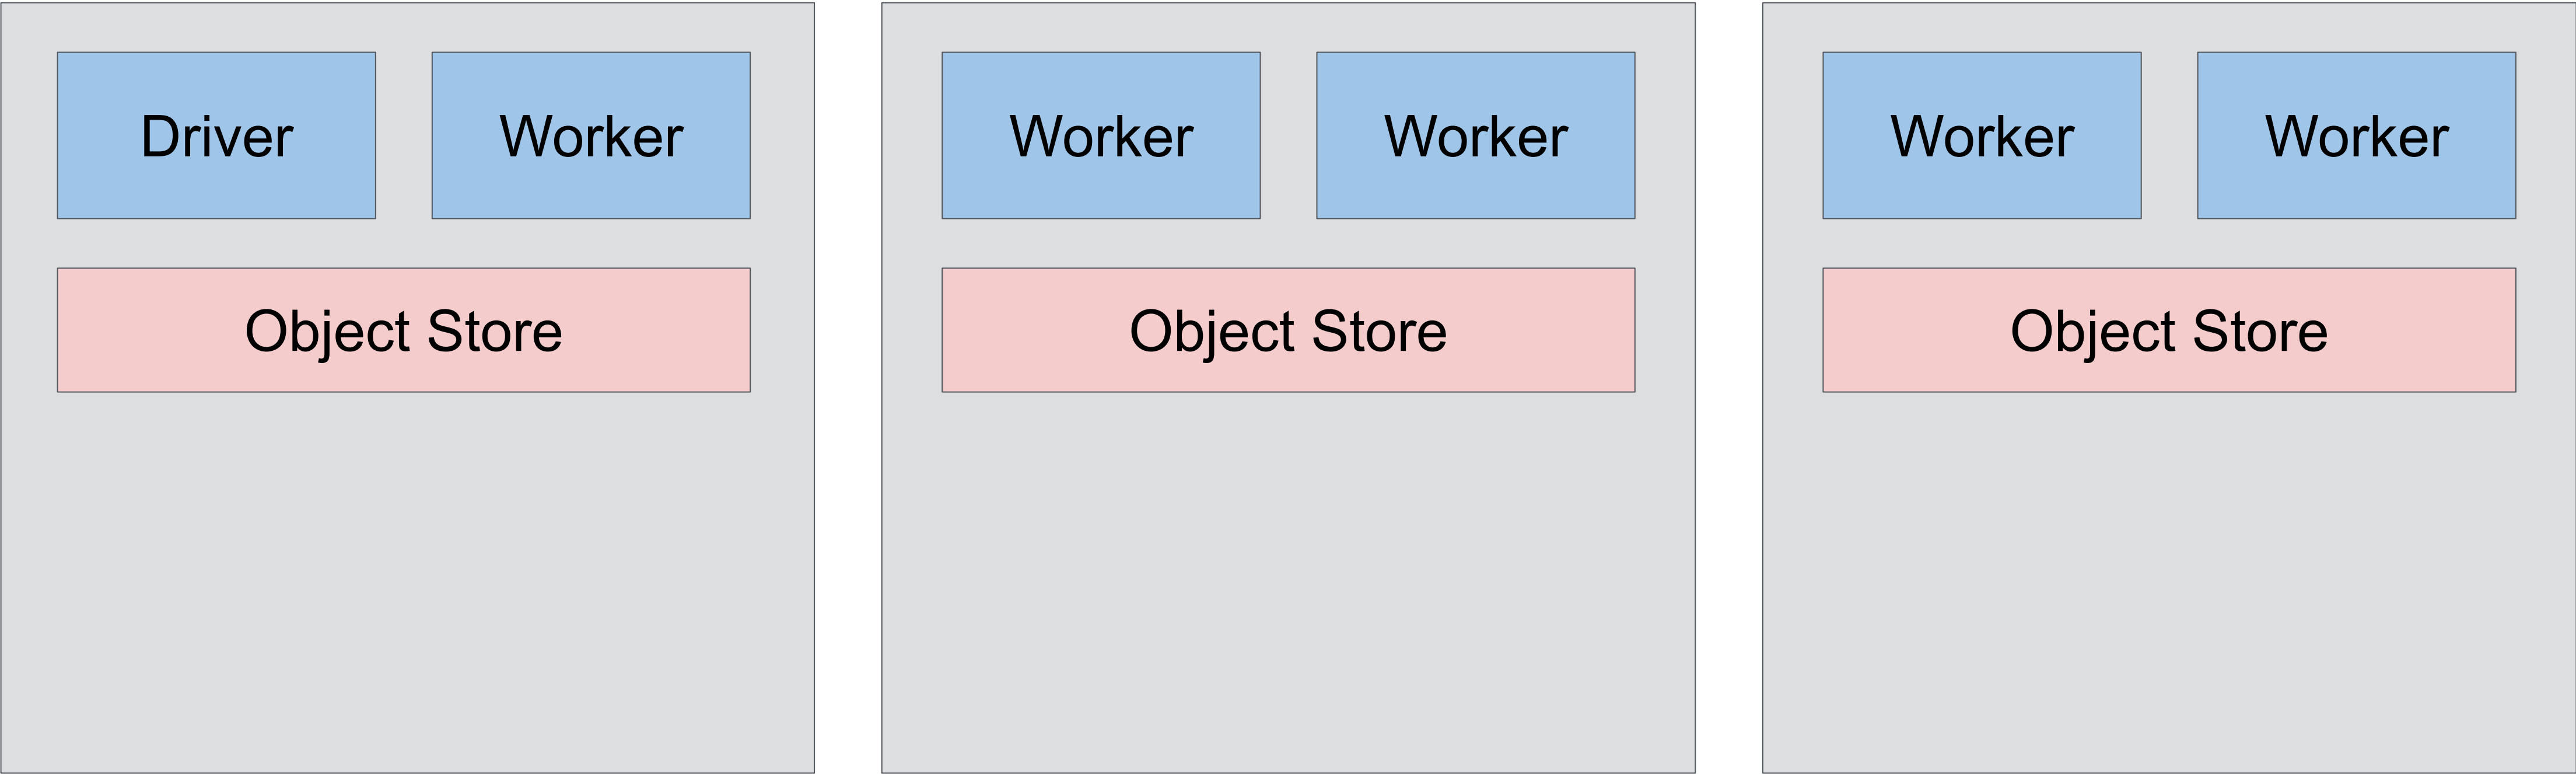
# Ray Architecture



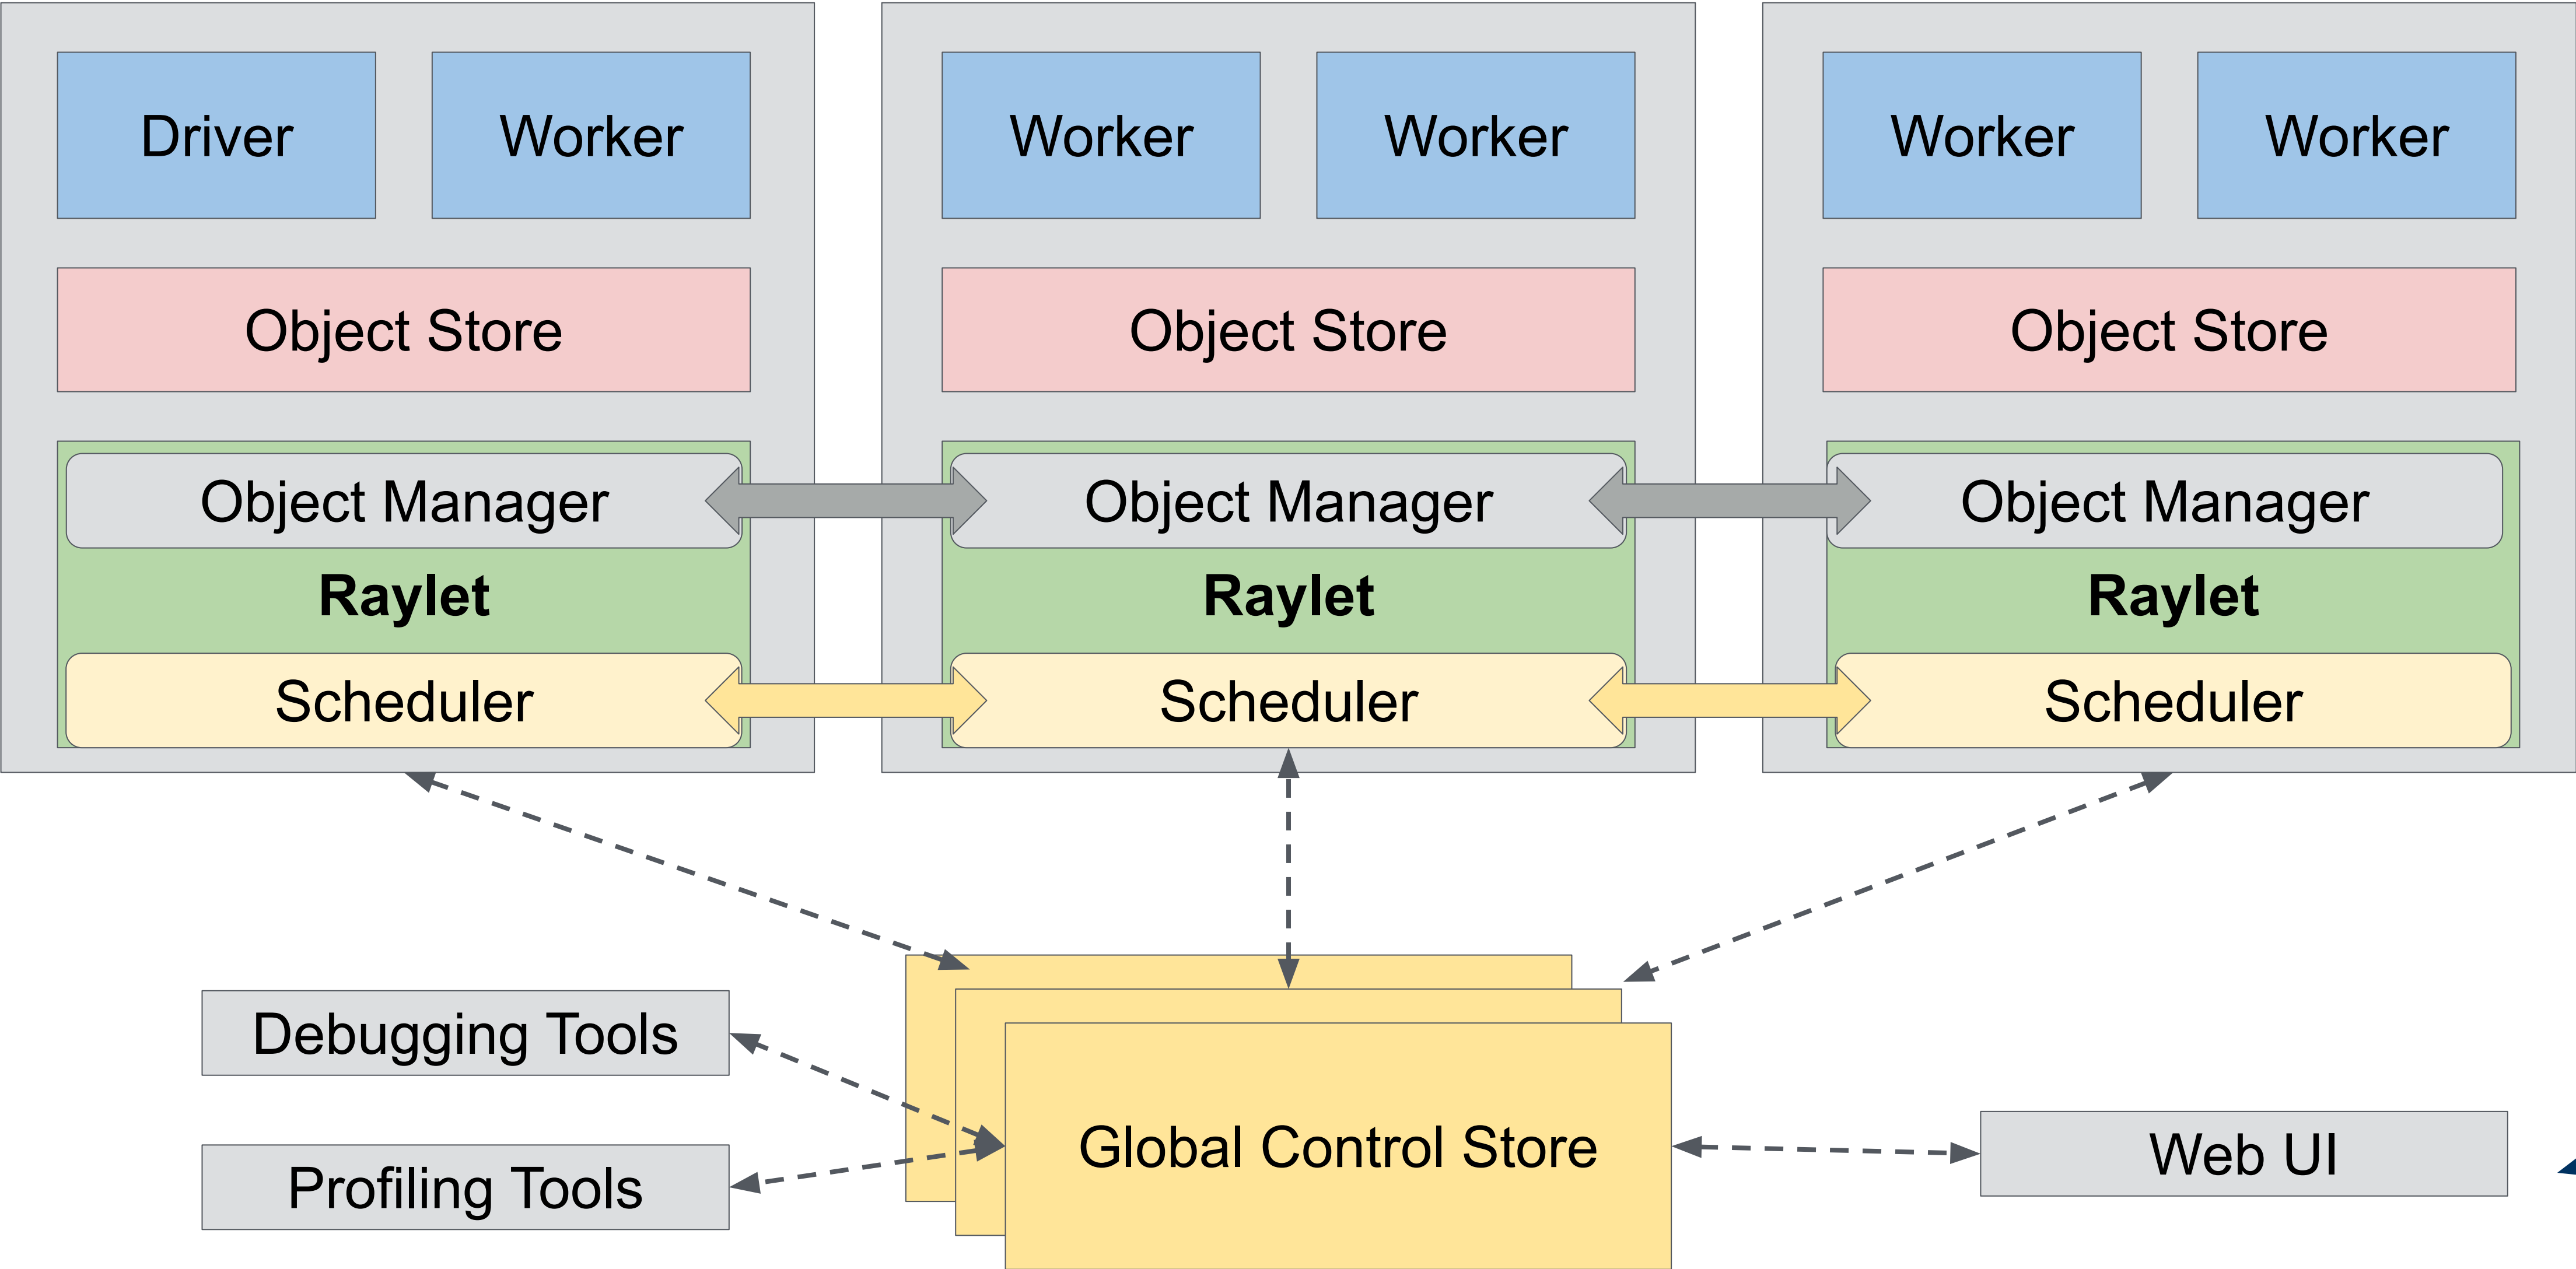
# Ray Architecture



# Ray Architecture



# Ray Architecture





# Exercises

Content available at [github.com/ray-project/tutorial/](https://github.com/ray-project/tutorial/).

## **Part 1: Ray as a unifying distributed system for machine learning**

**- exercises/ - 1, 2, (3 optional)**

Click on the “Binder” link and wait a couple minutes.

Binder not working? Backup:

- run locally with "pip install ray[rllib] tensorflow jupyterlab"





## **Distributed Hyperparameter Search on Ray**

# What is Tune?



Distributed  
Training

Model  
Serving

Streaming

Distributed RL

Data  
Processing

**Hyperparameter  
Search**

**Ray Libraries**

Tasks

Actors

**Ray API**


Dynamic Task Graphs

**Ray backend**

**Distributed System (Ray)**

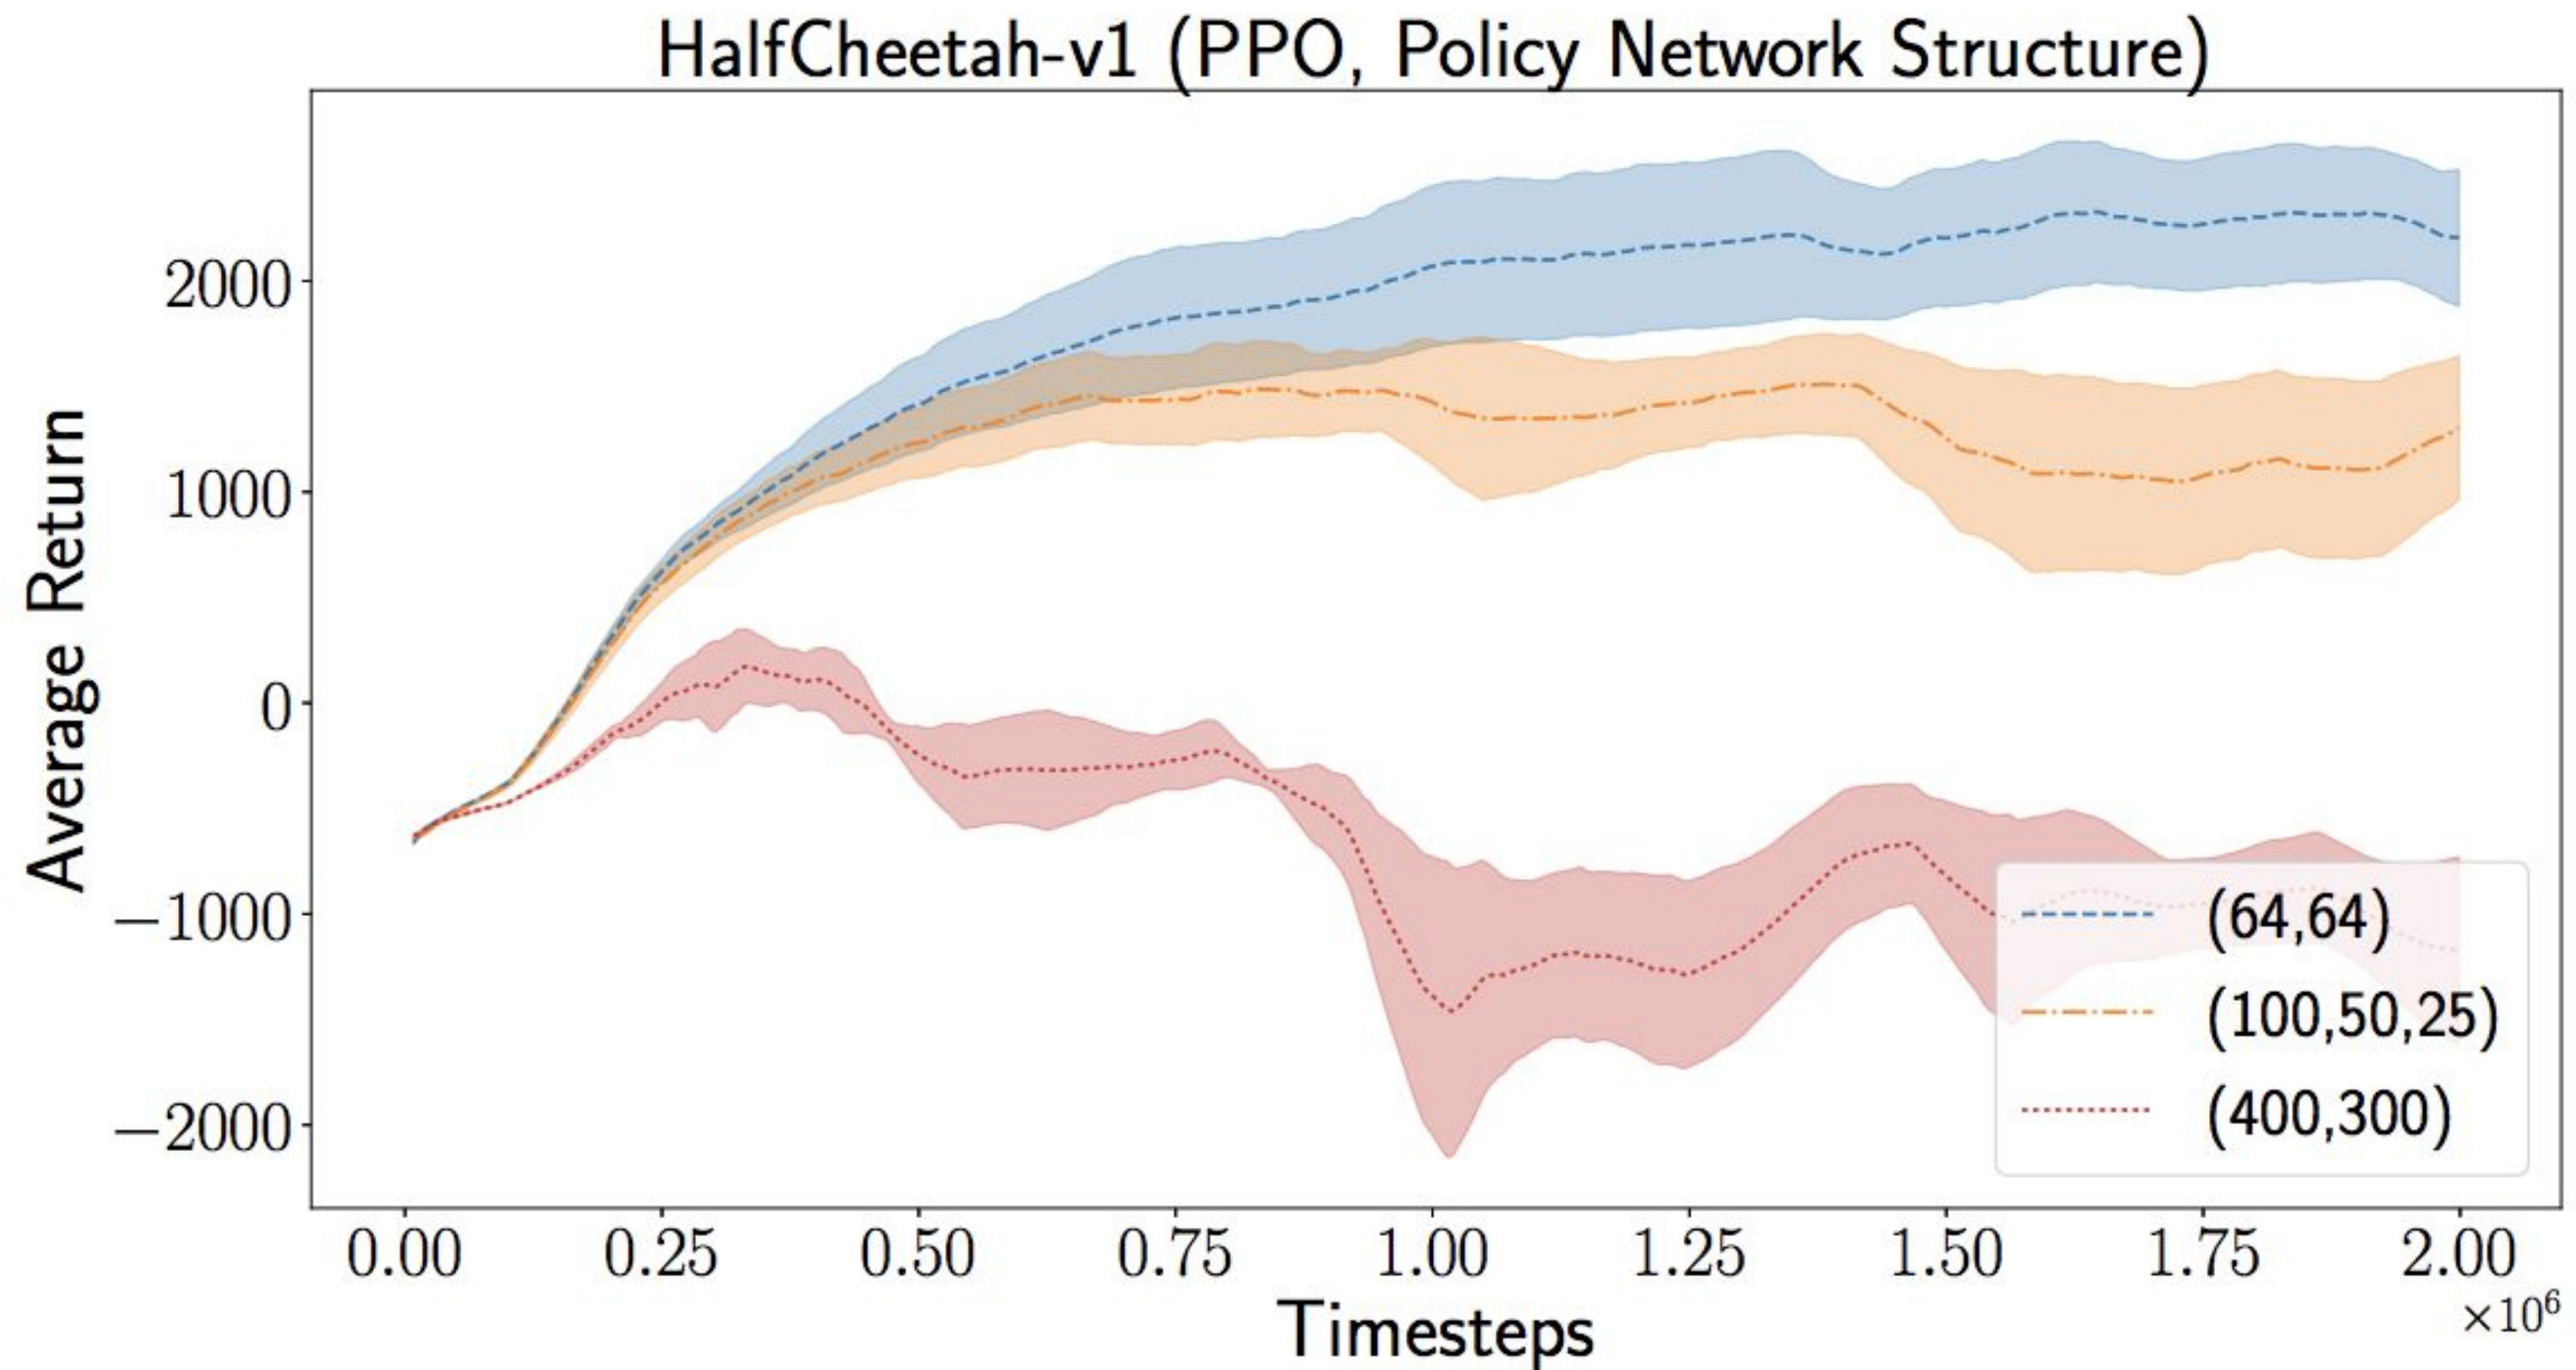
# Hyperparameters?







# Are hyperparameters actually that important?




# Why a framework for tuning hyperparameters?

We want the best  
model

Resources are  
expensive

Model training is  
time-consuming




# Tune is built with Deep Learning as a priority.

## Resource Aware Scheduling




## Seamless Distributed Execution



## Simple API for new algorithms

```
class TrialScheduler:  
    def on_result(self, trial, result): ...  
    def choose_trial_to_run(self): ...
```

## Framework Agnostic



[ray.readthedocs.io/en/latest/tune.html](https://ray.readthedocs.io/en/latest/tune.html)



# Tune is simple to use.

```
# Function-based API
def train():
    for _ in range(N):
        reporter(...)
```

```
# Class-based API
class MyModel(Trainable):
    def _setup(); def _train();
    def _save(); def _restore();
```

Two simple APIs  
for model training

## Tune API

Bayesian  
Optimization

Grid Search

...

HyperBand

Population  
Based Training

Pair **search algorithms**  
and **trial schedulers** to  
guide your distributed  
optimization.

# Exercises

Content available at [github.com/ray-project/tutorial/](https://github.com/ray-project/tutorial/).

## Part 2: Hyperparameter Tuning with Tune

- `tune_exercises/` - `Tutorial.ipynb`

Click on the “Binder” link and wait a couple minutes.

Binder not working? Backup:

- run locally with `"pip install ray[rllib] tensorflow jupyterlab"`



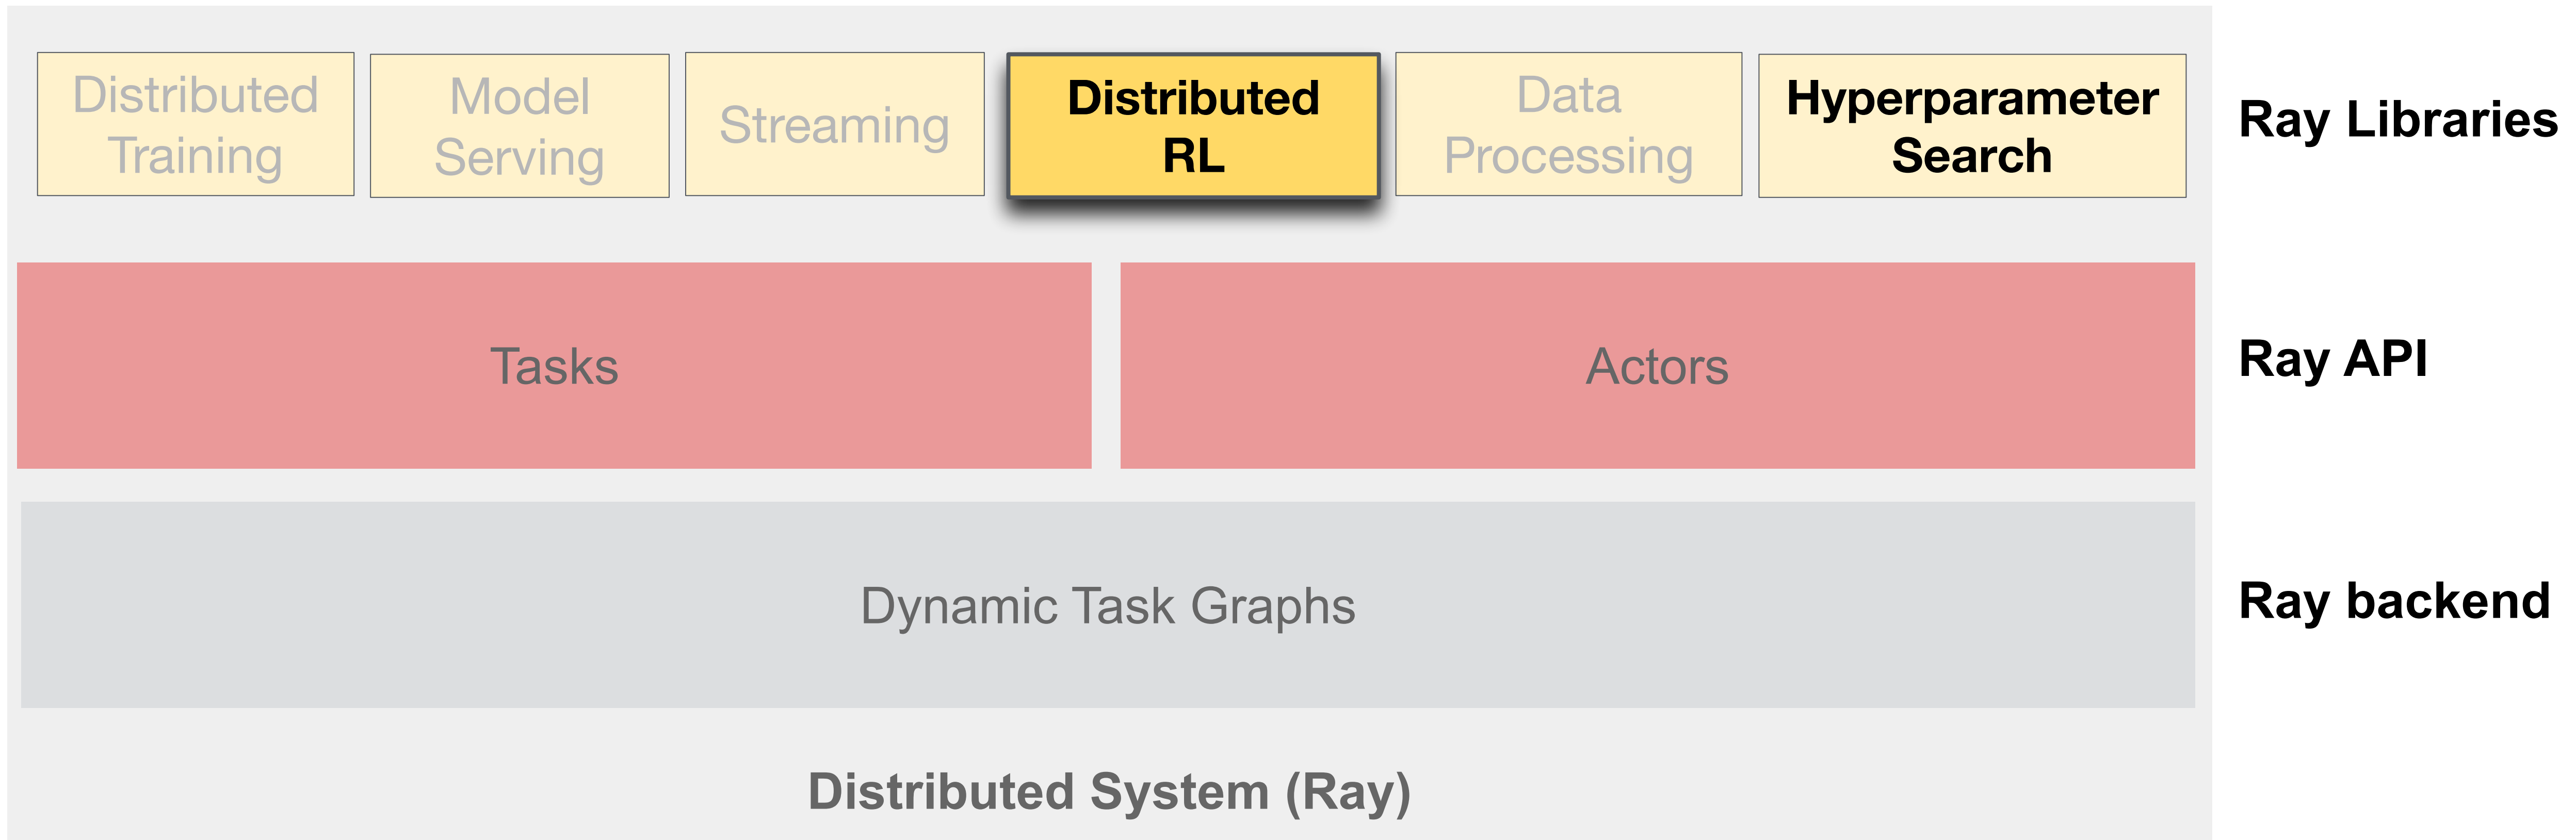


# RLlib


A scalable and unified library for reinforcement learning

<https://rllib.io>

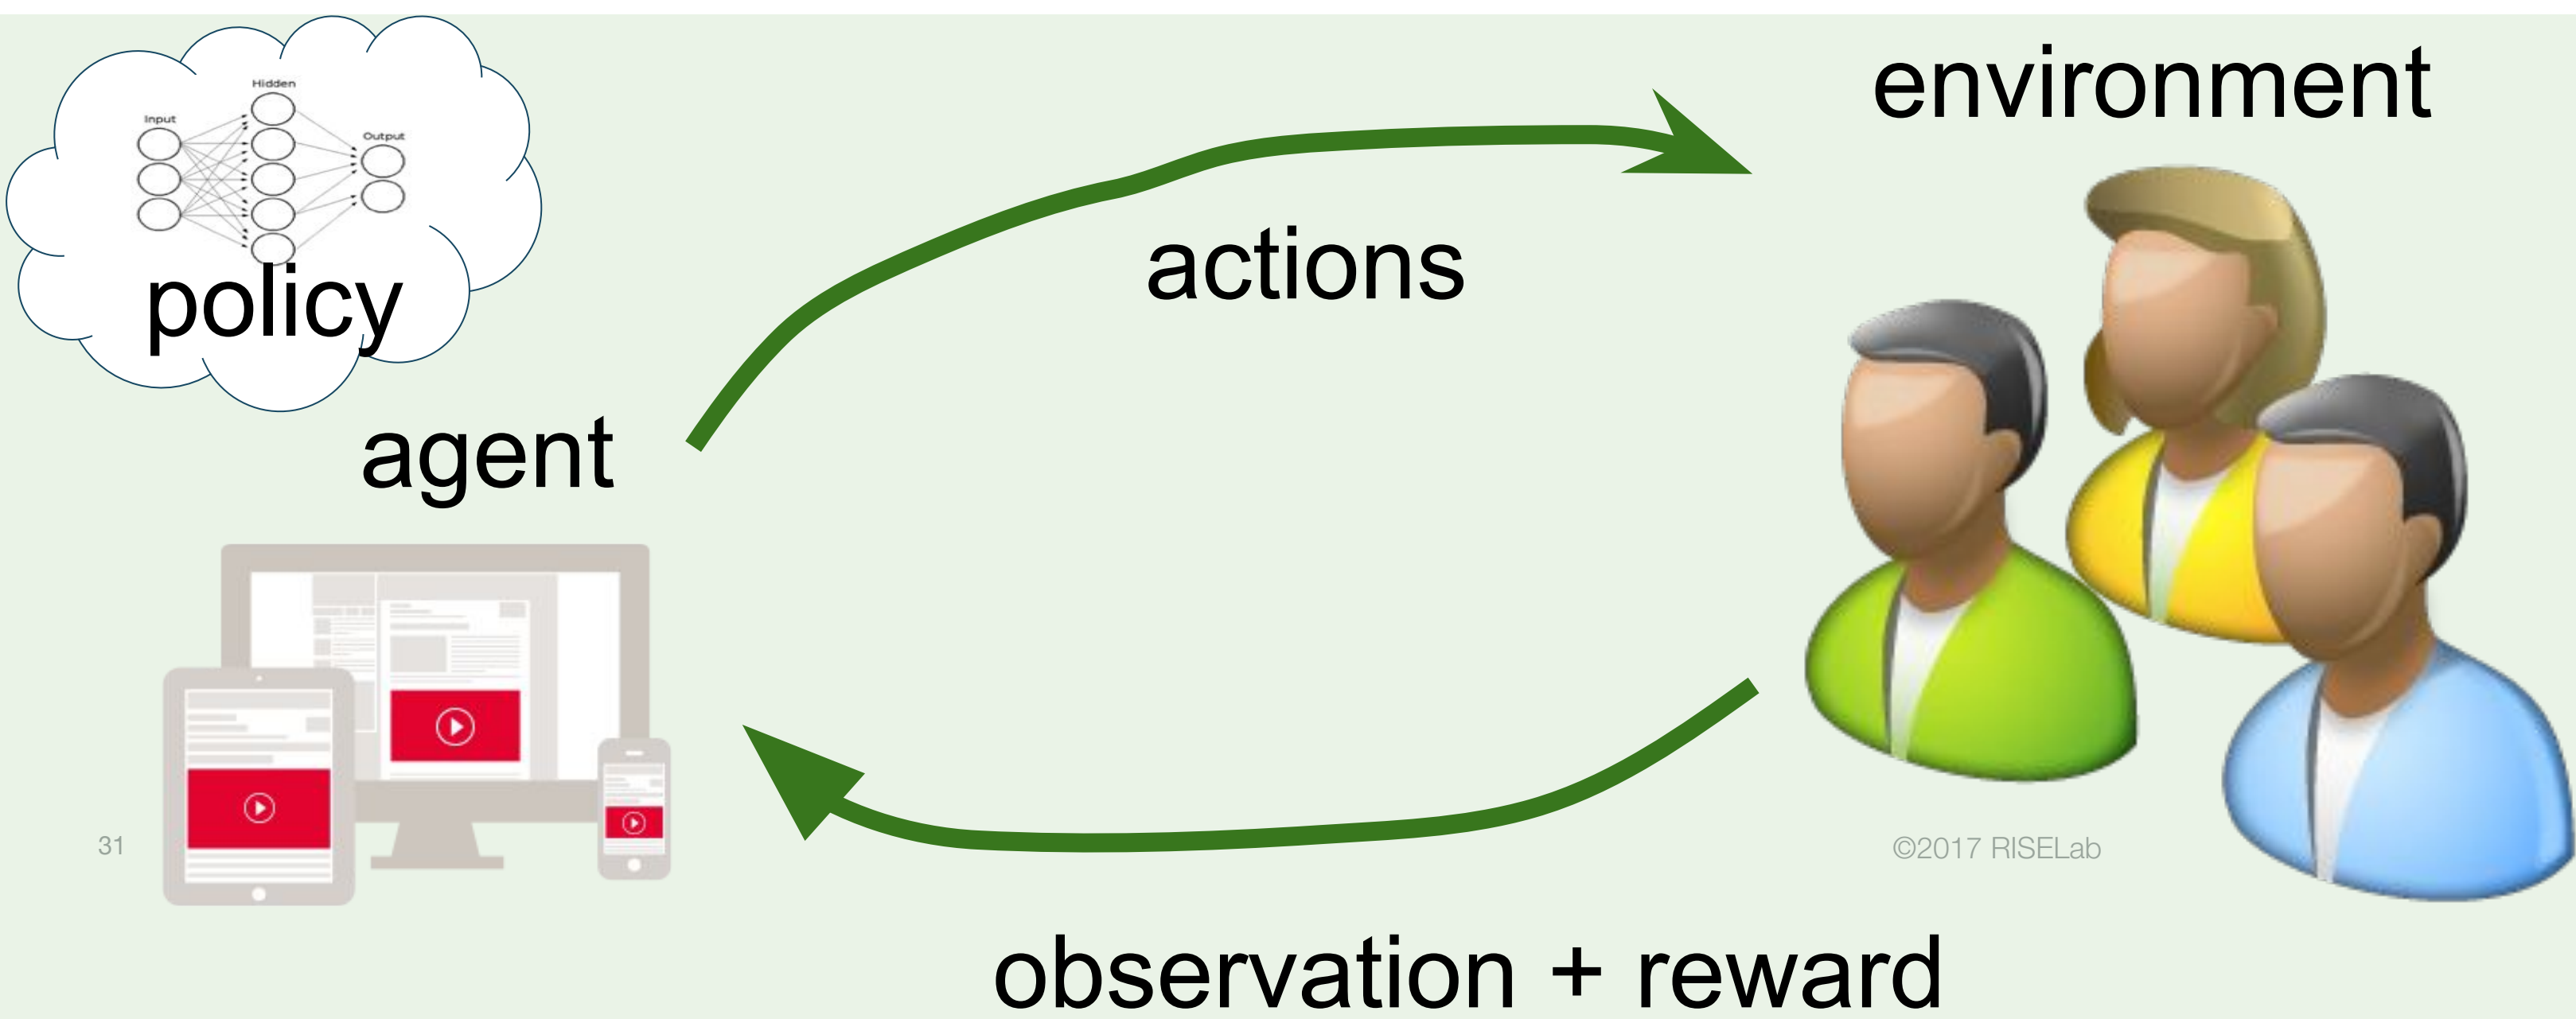
# What is RLlib?



# Background: What is reinforcement learning?



Supervised Learning



Reinforcement Learning



# Growing number of RL applications

Robotics


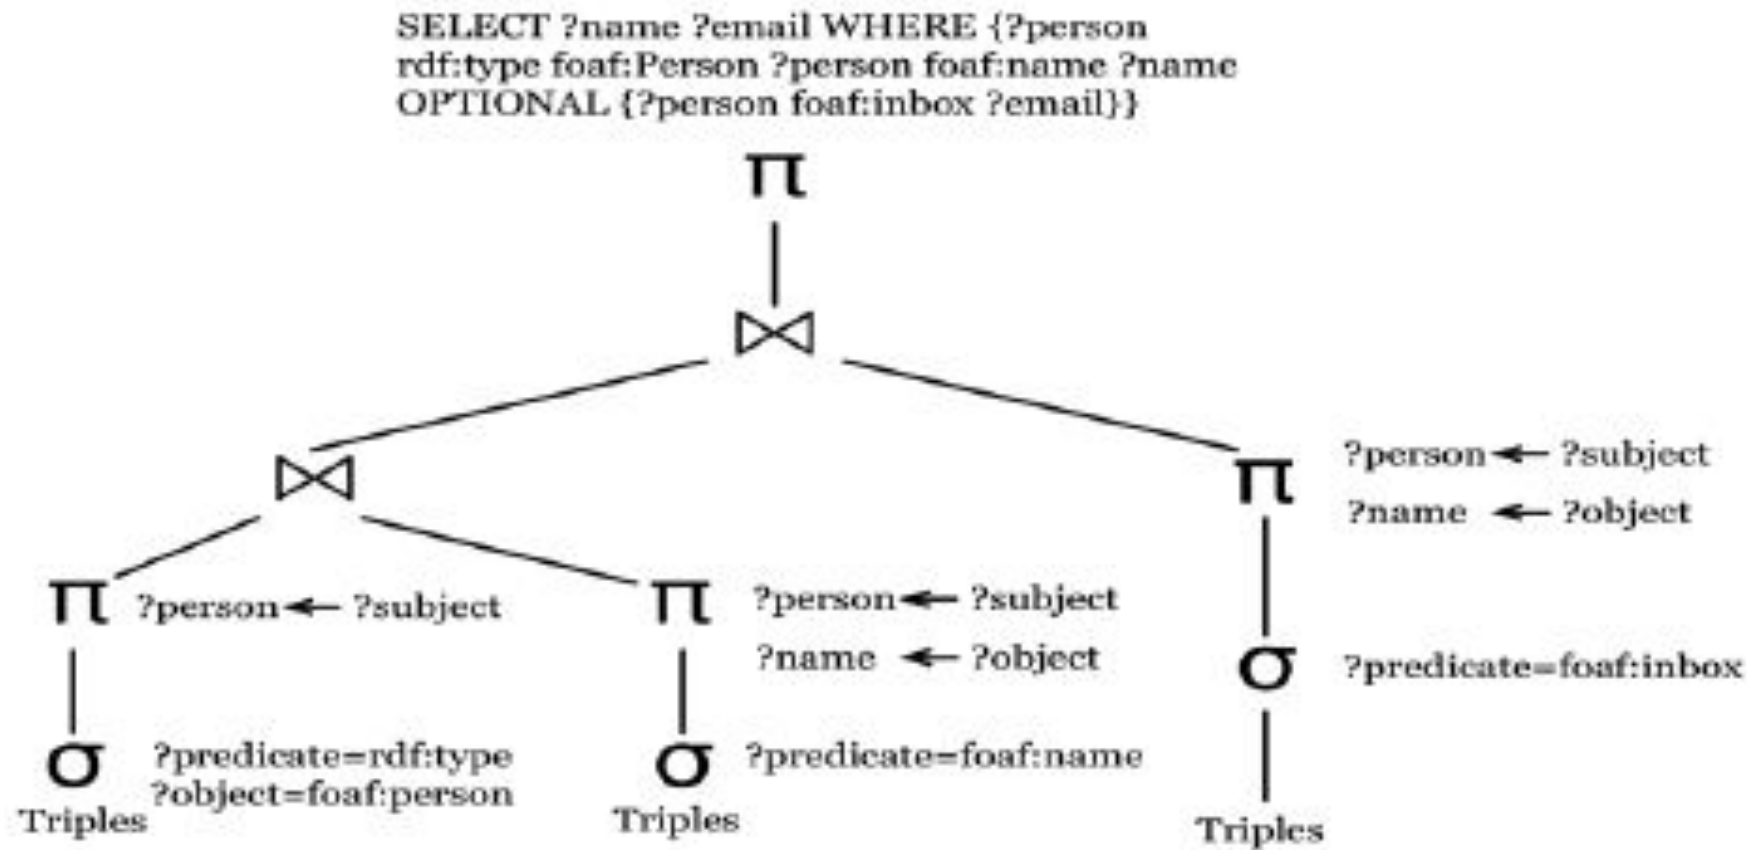

Industrial Control

Advertising

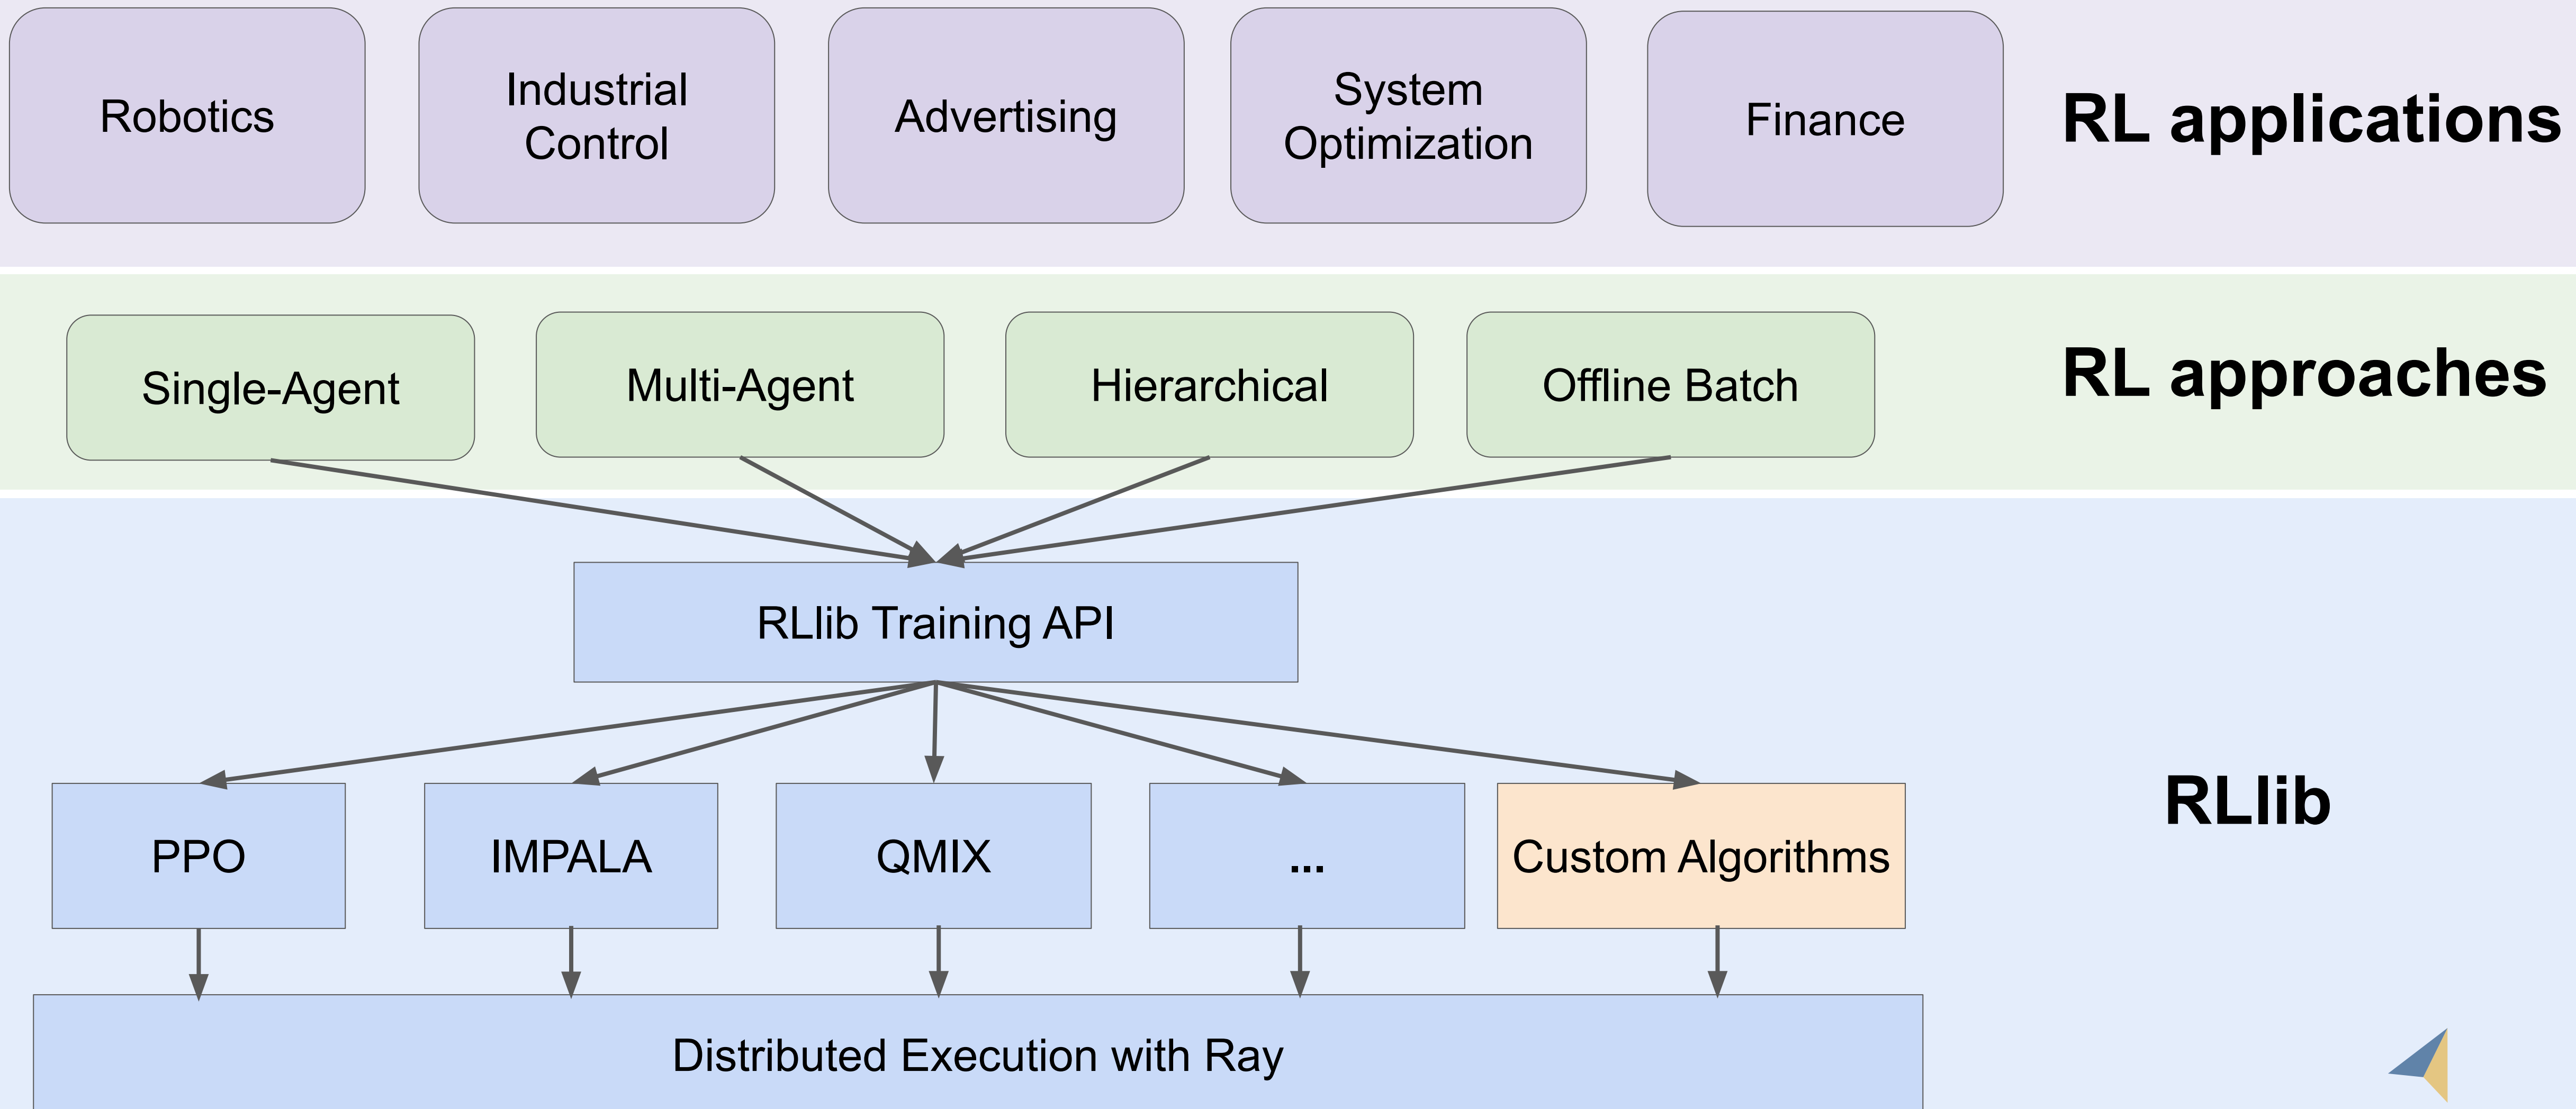
System Optimization

Finance

RL applications



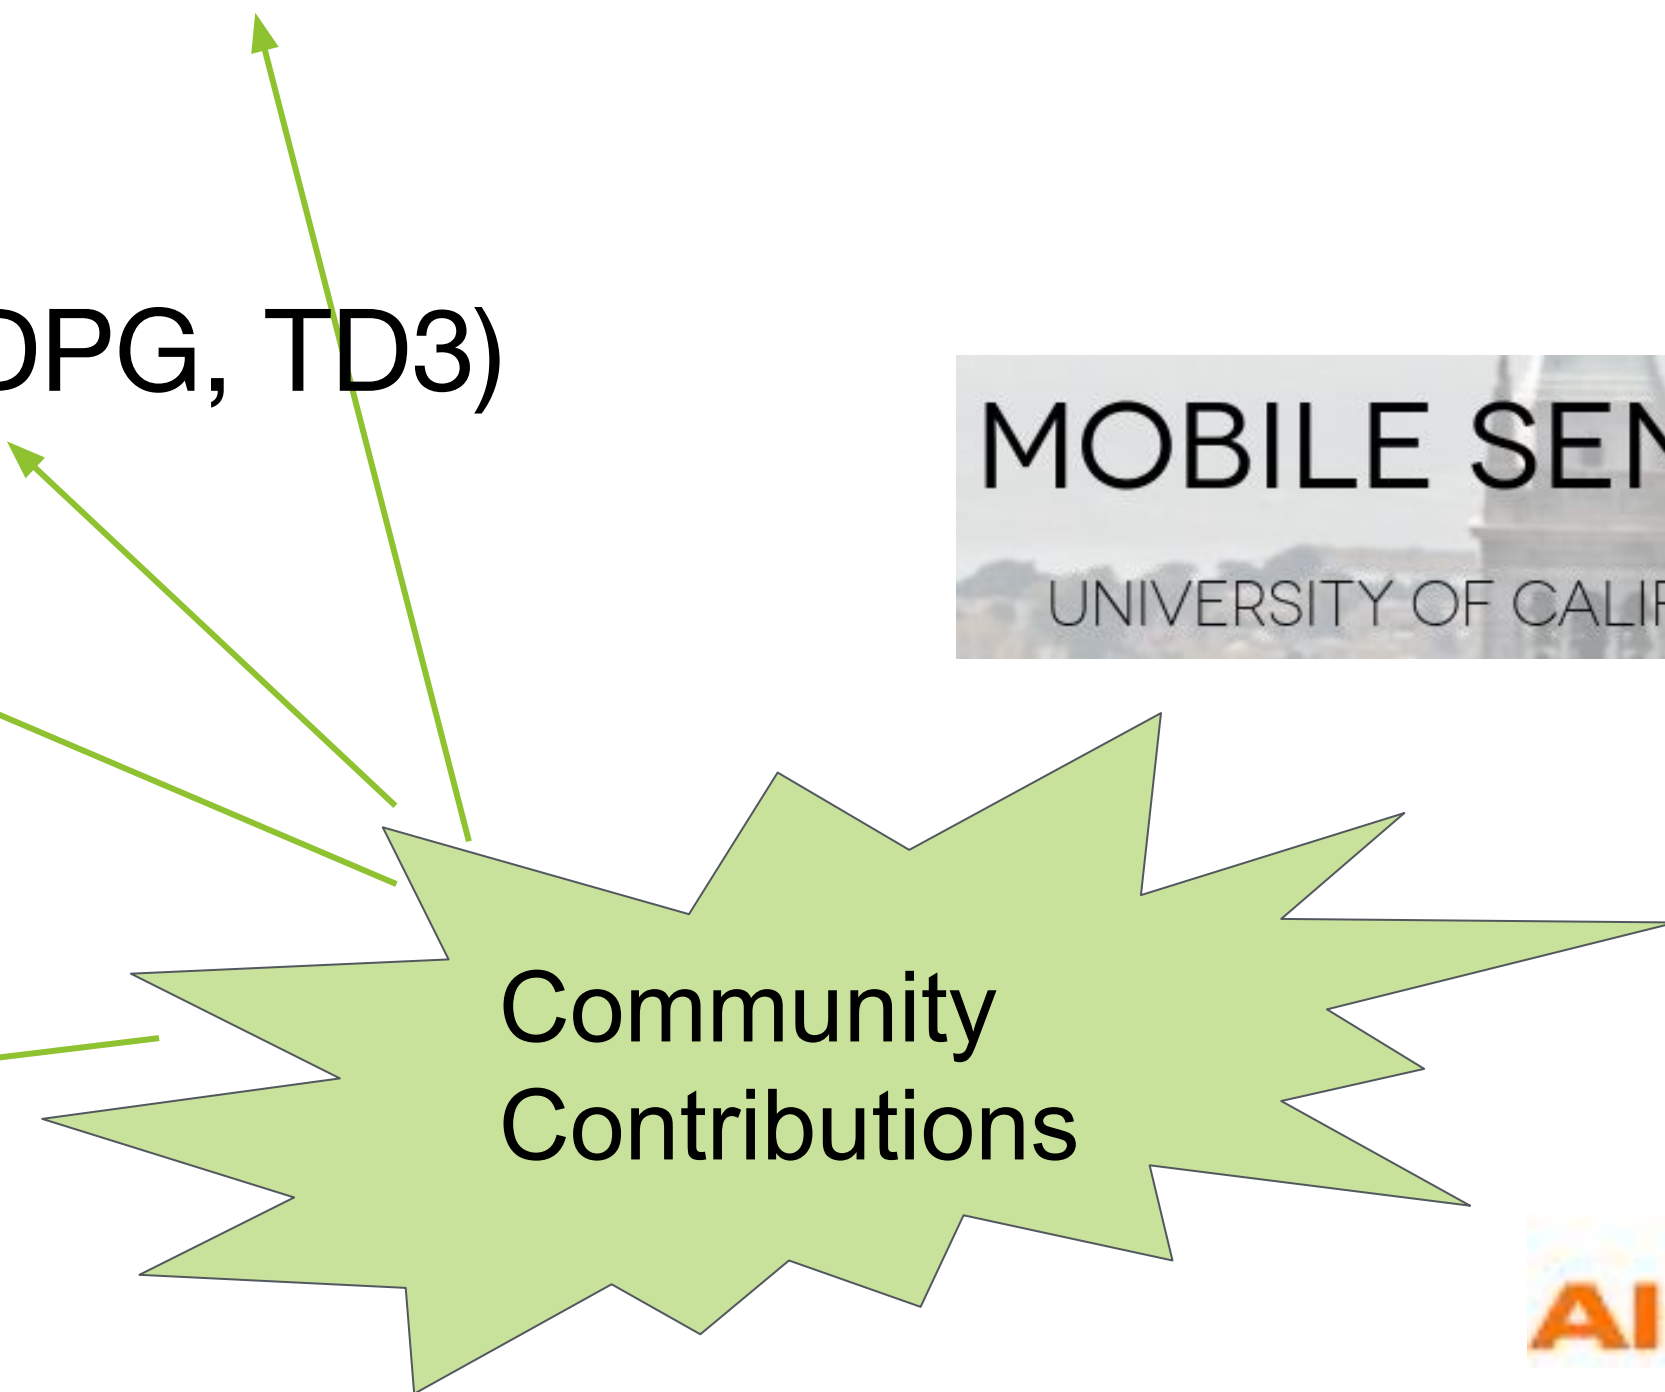
# A scalable, unified library for reinforcement learning






# Reference Algorithms

- **High-throughput architectures**
  - Distributed Prioritized Experience Replay (Ape-X)
  - Importance Weighted Actor-Learner Architecture (IMPALA)
- **Gradient-based**
  - Advantage Actor-Critic (A2C, A3C)
  - Deep Deterministic Policy Gradients (DDPG, TD3)
  - Deep Q Networks (DQN, Rainbow)
  - Policy Gradients
  - Proximal Policy Optimization (PPO)
- **Derivative-free**
  - Augmented Random Search (ARS)
  - Evolution Strategies

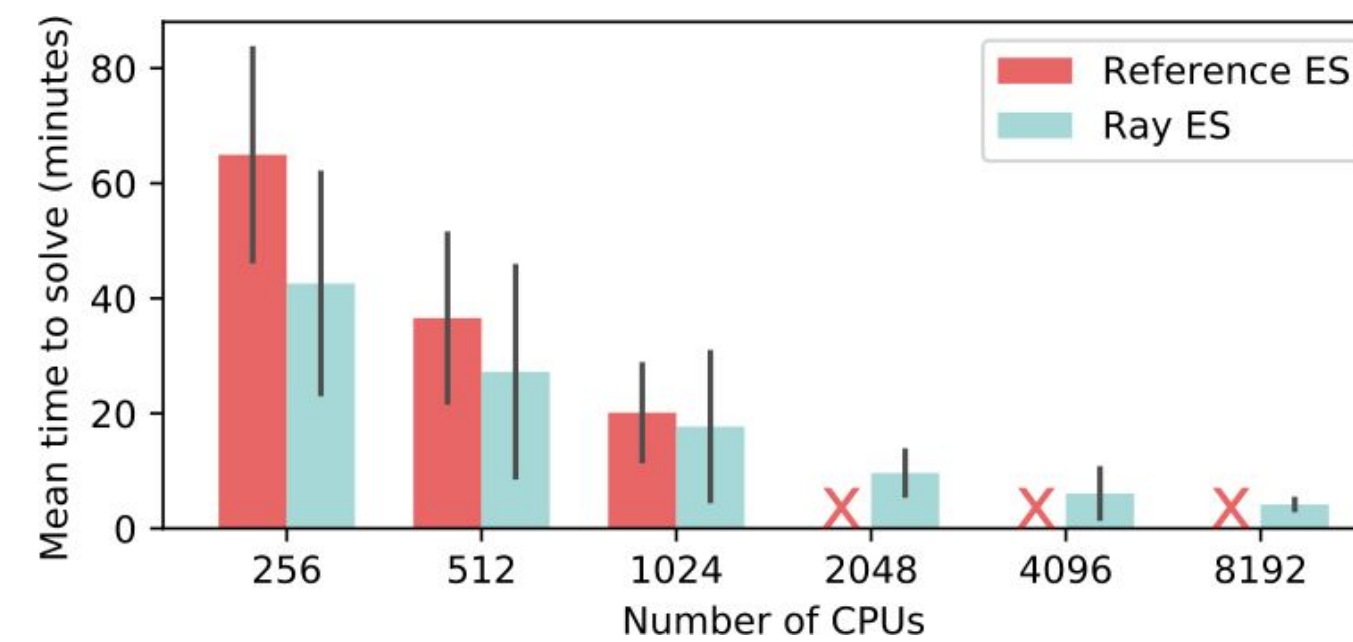


# Performance


## Distributed PPO (vs OpenMPI)



## Evolution Strategies (vs Redis-based)



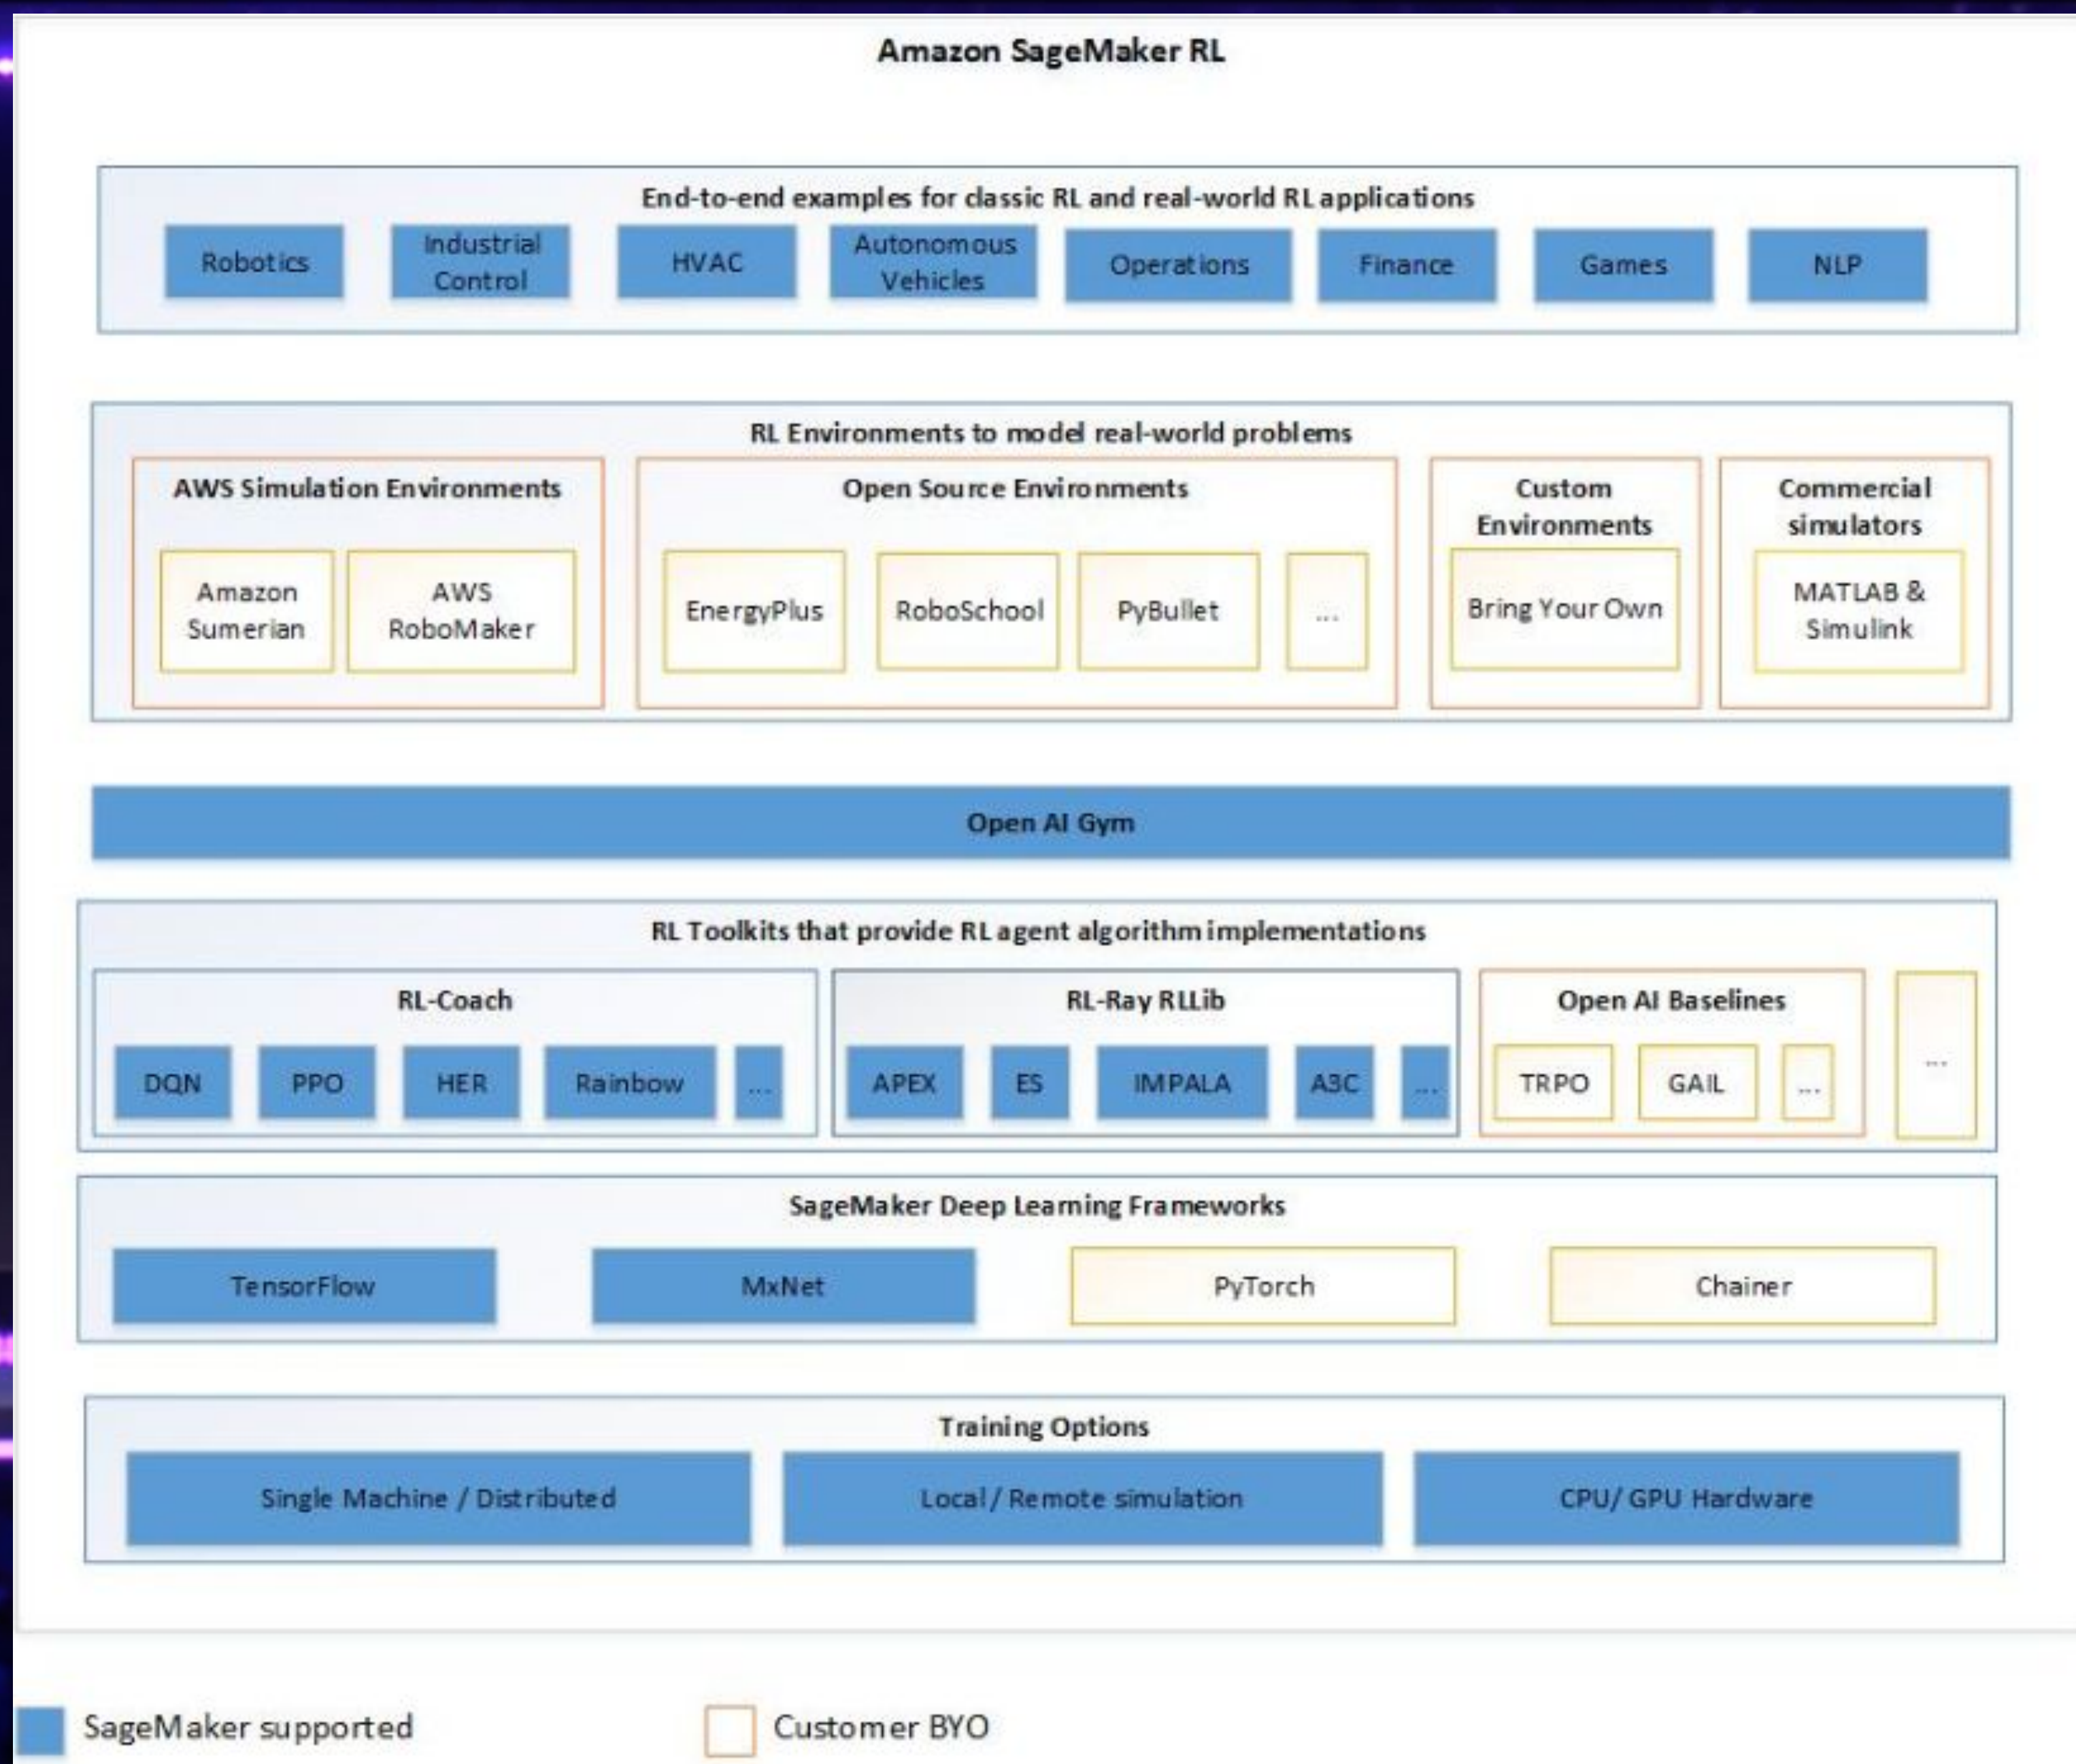
## Ape-X Distributed DQN, DDPG





# Amazon SageMaker RL

Reinforcement learning for every developer and data scientist





# Exercises

Content available at [github.com/ray-project/tutorial/](https://github.com/ray-project/tutorial/).

## Part 3: Applied Reinforcement Learning

- `rllib_exercises/` - `exercise01`, `exercise02`

Click on the “Binder” link and wait a couple minutes.

Binder not working? Backup:

- run locally with `"pip install ray[rllib] tensorflow jupyterlab"`





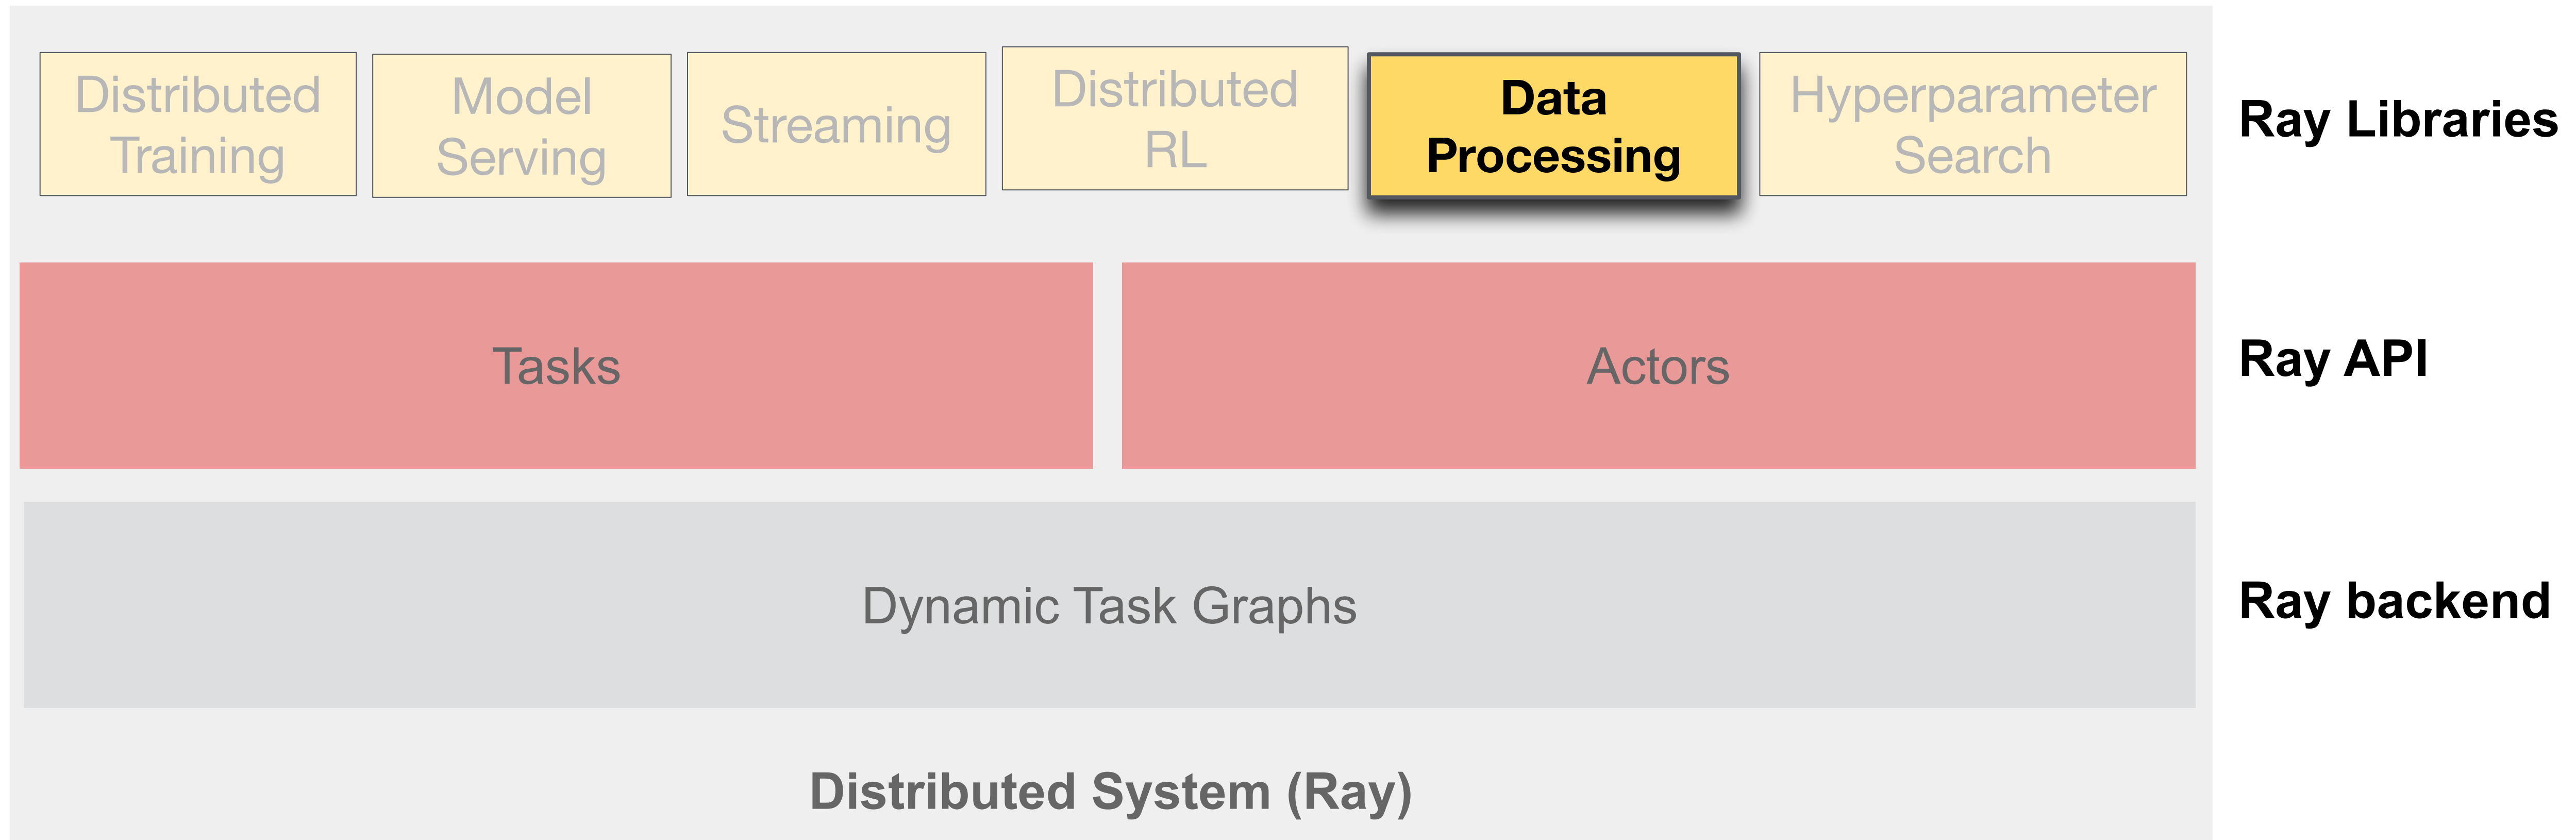
**Accelerate your Pandas workflows by  
changing a single line of code**



[ray.readthedocs.io/en/latest/tune.html](https://ray.readthedocs.io/en/latest/tune.html)

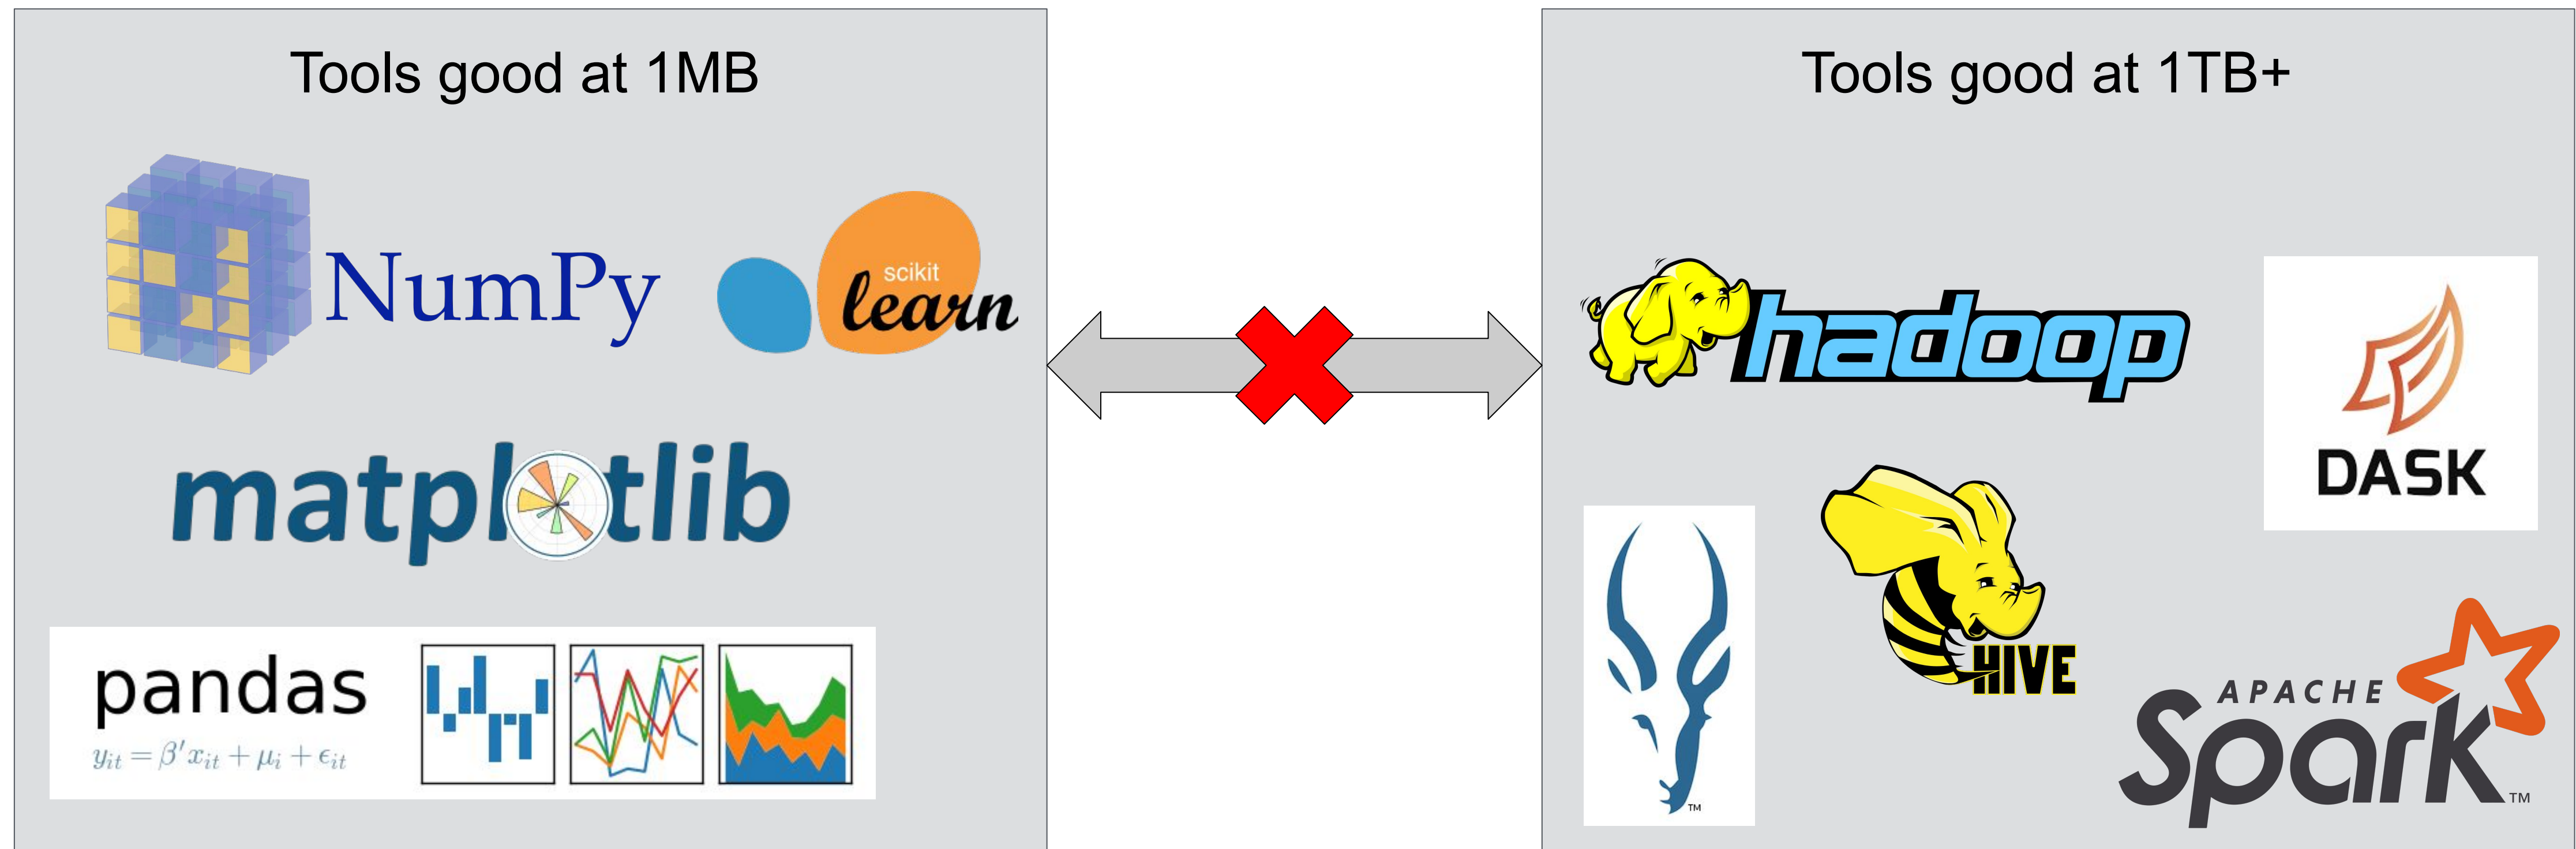


# What is Modin?



# Modin: Pandas on Ray

Accelerate your pandas workloads by changing one line of code



# Modin: Pandas on Ray

Accelerate your pandas workloads by changing one line of code

To use Modin, replace the pandas import:

```
# import pandas as pd  
import modin.pandas as pd
```

## Installation

Modin can be installed from PyPI:

```
pip install modin
```

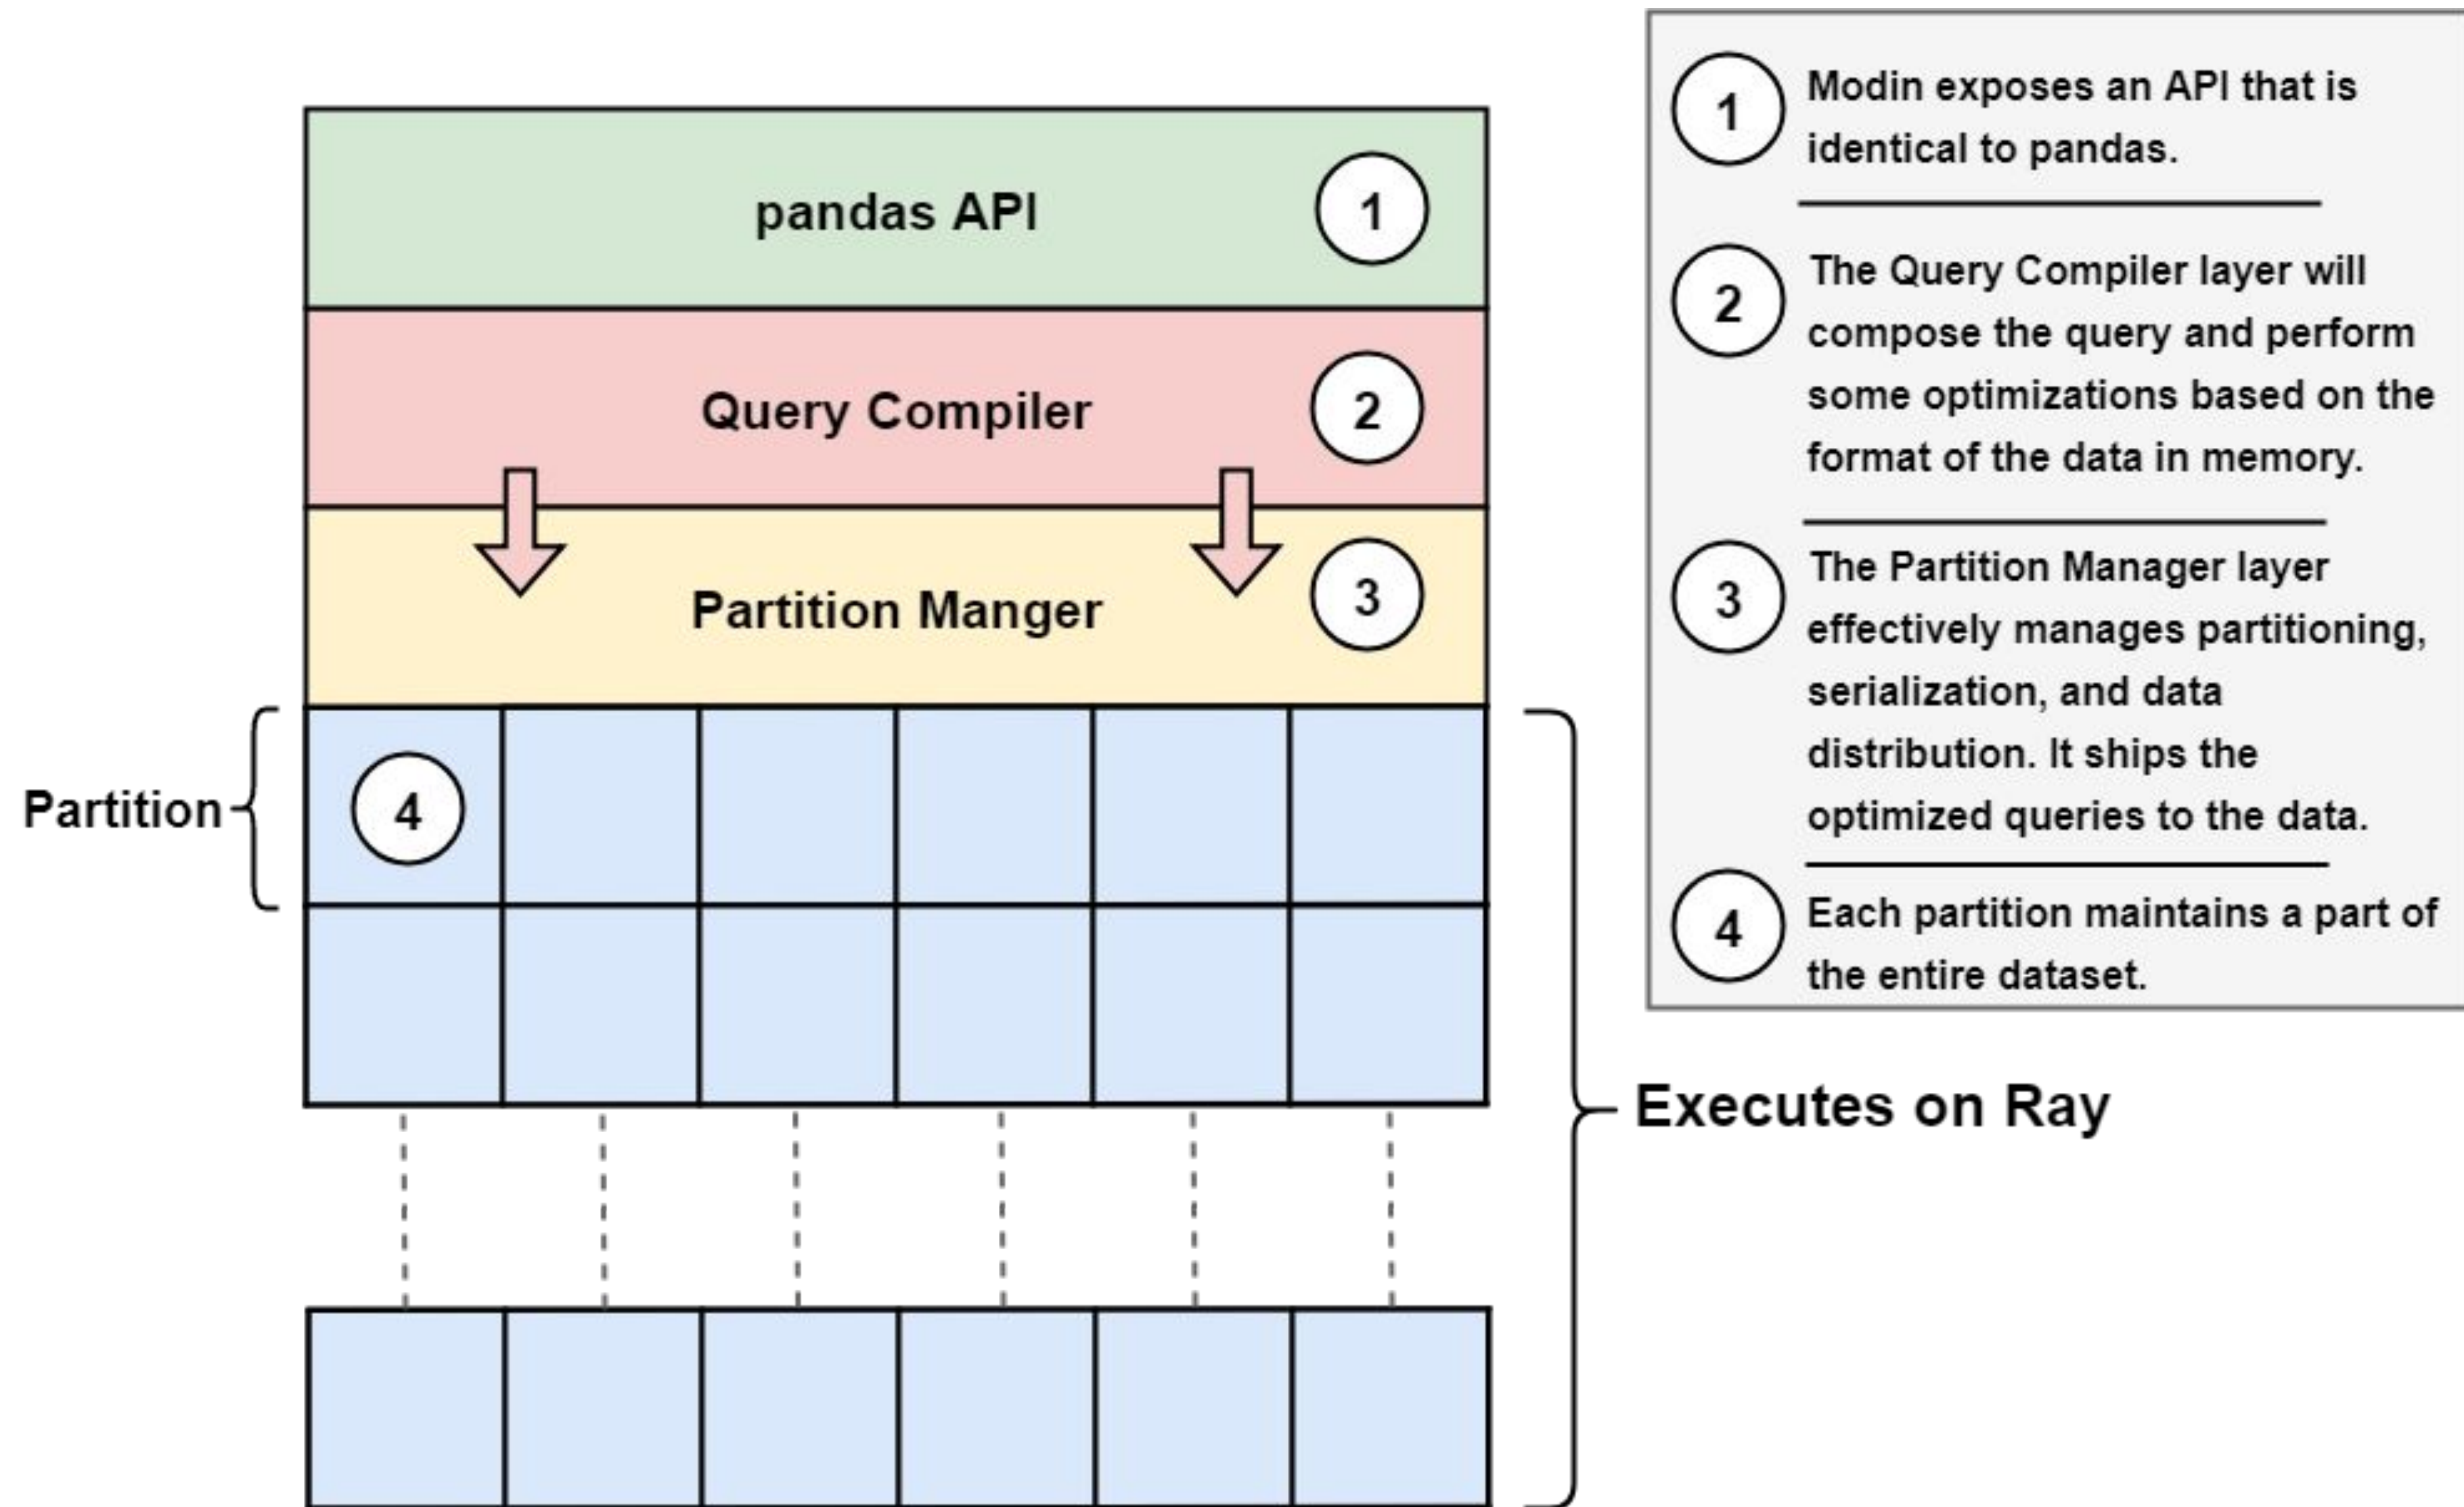


# Why Modin?

- Faster pandas, even on your laptop
  - Up to 4x speed improvement over pandas on 4 physical cores
- Cluster support -- experimental!
- A DataFrame library aimed at bridging the gap between MB-scale and TB-scale data






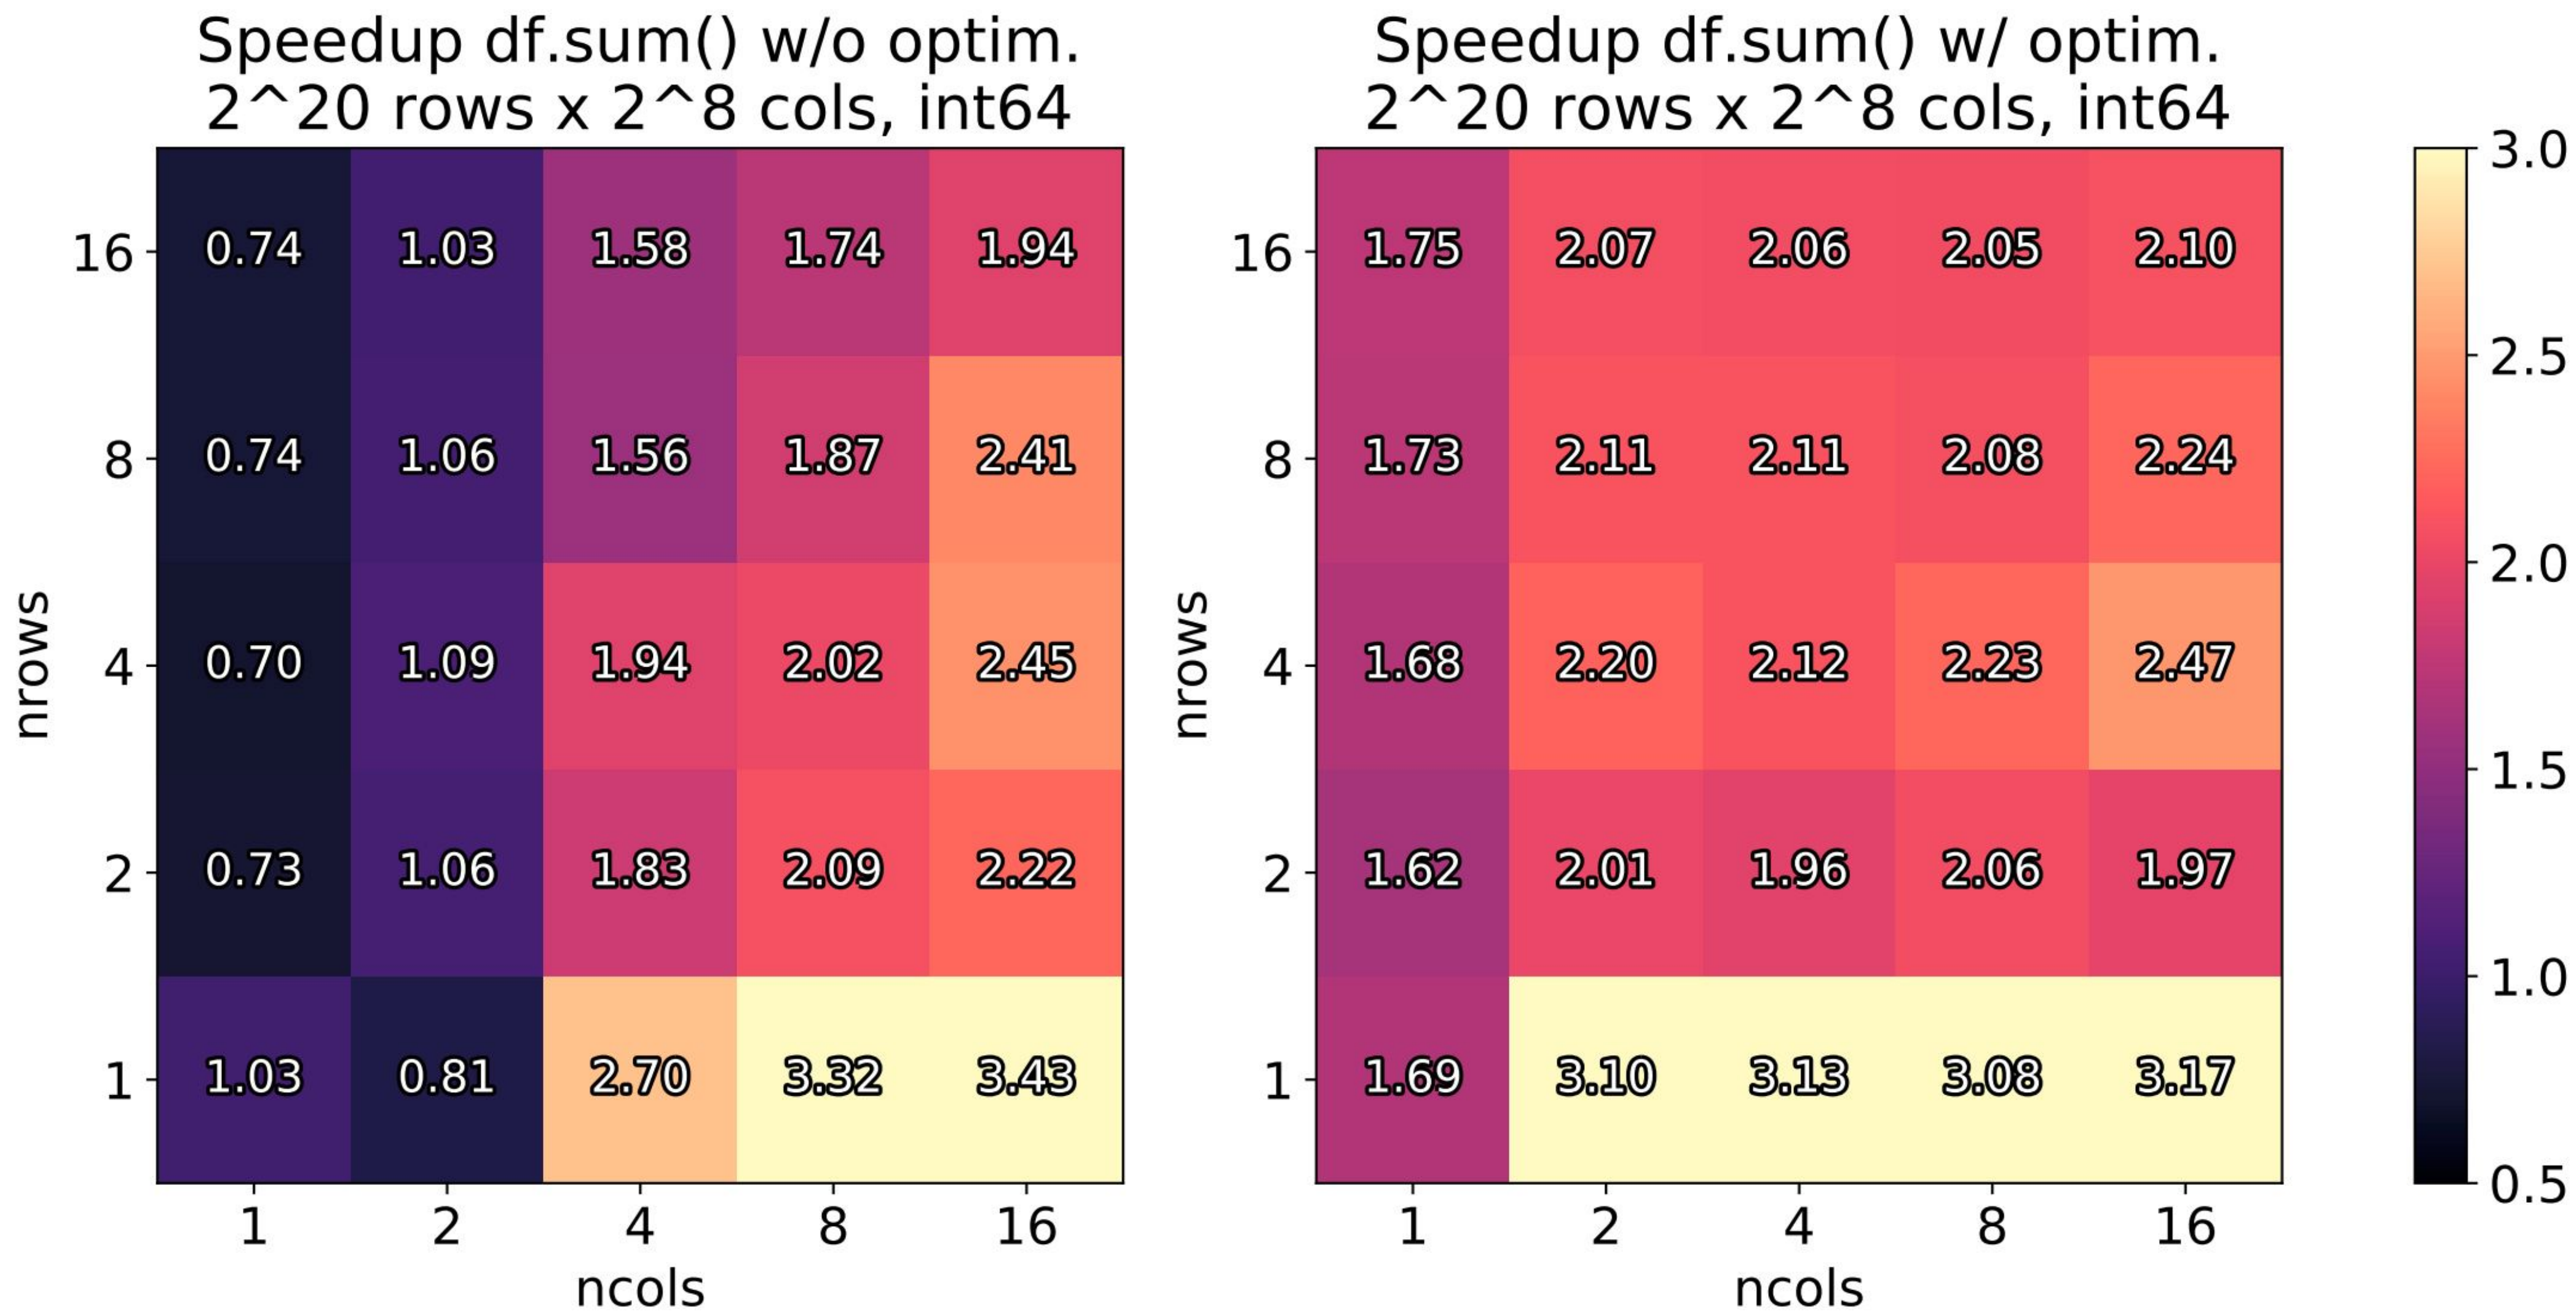




# Performance



# Performance



# Conclusion

- Ray is an open source project for distributed computing
- **special-purpose** distributed systems -> **general-purpose** distributed system
- Support for the full ML lifecycle (data collection, training, simulation, serving)



[github.com/ray-project/ray](https://github.com/ray-project/ray)

©2017 RISELab

